



BRNO UNIVERSITY OF TECHNOLOGY

VYSOKÉ UČENÍ TECHNICKÉ V BRNĚ

FACULTY OF INFORMATION TECHNOLOGY

FAKULTA INFORMAČNÍCH TECHNOLOGIÍ

DEPARTMENT OF INTELLIGENT SYSTEMS

ÚSTAV INTELIGENTNÍCH SYSTÉMŮ

ABSTRACTION OF STATE LANGUAGES IN AUTOMATA ALGORITHMS

ABSTRAKCE JAZYKŮ STAVŮ V AUTOMATOVÝCH ALGORITMECH

BACHELOR'S THESIS

BAKALÁŘSKÁ PRÁCE

AUTHOR

AUTOR PRÁCE

DAVID CHOCHOLATÝ

SUPERVISOR

VEDOUCÍ PRÁCE

Doc. Mgr. LUKÁŠ HOLÍK, Ph.D.

BRNO 2022

Abstract

Finite automata theory is a well-known model of computational theory. We use finite automata daily in many situations without even knowing it. In this thesis, we will focus our attention on different heuristics for optimizing several typical problems connected to finite automata. We are interested in possibilities of abstraction of automata state languages to be used in various automata algorithms. We explore especially the computation of intersection of automata product construction and its emptiness test, which are required in modern computation technologies, but requires a lot of computational time and generates unnecessary state space. For this reason, we use various optimization approaches for these operations such as length abstraction, Parikh image computation or mintermization for solving mentioned problems and optimizing the product construction and its emptiness test using knowledge gained from abstractions of automata state languages.

Abstrakt

Teorie konečných automatů je dobře známou oblastí teorie algoritmů. Konečné automaty běžně denně používáme v mnohých situacích, aniž bychom si to uvědomovali. V této práci se zaměříme na rozdílné heuristiky pro optimalizaci několika typických problémů spjatých s konečnými automaty. Zaujímáme se převážně o možnosti abstrakce jazyků stavů pro jejich použití v automatových algoritmech. Studujeme převážně konstrukci průniku konečných automatů a jeho test prázdnoty, kterých je často využíváno v moderních výpočetních technologiích, ale které vyžadují značný výpočetní čas a produkují rozsáhlý stavový prostor. Z tohoto důvodu využíváme přístupů optimalizace těchto operací jako jsou délková abstrakce, Parikovsky obrazy nebo mintermizace pro řešení zmíněných problémů a optimalizaci konstrukce průniku a jeho testu prázdnoty s využitím znalostí získaných z abstrakcí jazyků automatových stavů.

Keywords

Finite Automata, State Languages Abstraction, Product Construction, Emptiness Test, Intersection Computation Optimization, State Space Reduction, SMT Solving, Length Abstraction, Parikh Image, Mintermization

Klíčová slova

Konečné automaty, Abstrakce jazyků stavů, Konstrukce produktu, Test prázdnoty, Optimalizace výpočtu průniku, Redukce stavového prostoru, SMT výpočty, Délková abstrakce, Parikovsky obrazy, Mintermizace

Reference

CHOCHOLATÝ, David. *Abstraction of State Languages in Automata Algorithms*. Brno, 2022. Bachelor's thesis. Brno University of Technology, Faculty of Information Technology. Supervisor Doc. Mgr. Lukáš Holík, Ph.D.

Abstraction of State Languages in Automata Algorithms

Declaration

I hereby declare that this Bachelor's thesis was prepared as an original work by the author under the supervision of Doc. Mgr. Lukáš Holík, Ph.D. I have listed all the literary sources, publications and other sources, which were used during the preparation of this thesis.

.....
David Chocholatý
March 26, 2022

Acknowledgements

I would like to thank my supervisor, Doc. Mgr. Lukáš Holík, Ph.D., who has provided essential and necessary information about the topic, outlined possible solutions and answered every question I have had throughout the whole process.

Contents

1	Introduction	2
2	Preliminaries	5
3	State Language Optimizations	8
3.1	State Language Optimization with Length Abstraction	8
3.1.1	Length Abstraction Represented by Lasso Automata	9
3.1.2	Single Lasso Automaton for Each Original Automaton	11
3.2	Product Construction with Length Abstraction Optimization	11
3.2.1	Optimization with Skipping Satisfiable States	15
3.2.2	Resolving Length Abstraction Satisfiability without SMT Solver . .	17
3.3	Product Construction Optimization with Parikh Image Computation	17
3.3.1	Parikh Image	19
3.3.2	Optimization Algorithm Using Parikh Images	22
3.3.3	Optimization with Incremental SMT Solving	24
3.4	Optimization with SMT Solver Timeout	26
3.5	Combined Approach of Parikh Image Optimization Enhanced by Length Abstraction	26
3.6	Abstraction of State Language with Mintermization	27
4	Experiments and Results	31
5	Conclusion	35
	Bibliography	36

Chapter 1

Introduction

Automata theory is a well-known field of study. Automata appear in our mundane lives every day. It is consequently important to study automata further and learn to operate with them, too. Automata are commonly used in mathematics and computation theory in general (e.g., in model checking [18] or string solving and analysis [16]). Of course, their uses in the field of logic are not to be forgotten (e.g., WS1S [9, 10]).

Our goal in this thesis is to explore possibilities of using various abstractions of automata languages in optimization of automata algorithms. We will study different approaches to abstraction of languages of states. We start with abstracting languages of states to sets of possible word lengths and to Parikh images, represented as semi-linear sets, and exploring options of using them to optimize the construction of synchronous product of automata by pruning pairs of states with incompatible abstractions. We continue towards optimization of these techniques. Furthermore, we consider the use of mintermization and other approaches to further improve these methods.

Finite automata are easy to understand, but as soon as you start adding new states, transitions and additional accepting states, the finite automaton will get extensively larger and harder to work with. Our focus is on several specific, often used operations over the finite automata, which are usually simple, but take lots of computational time and generate vast state space as a result. One of such operations over the automata is construction of their intersection. We use product construction and its emptiness test as a benchmarking operation in our experiments. Our suggested optimization methods are tested on the product construction algorithm, but are generally usable on many other typical automata algorithms. Consequently, even if our approaches to the optimization problems are introduced on product construction algorithms, our discoveries have wider impact and are in some form applicable on other automata operations, e.g., complement generation etc.

The intersection of two or more automata (the synchronous product construction) is often used in mathematics and logic. It is an extensively used operation where one follows the transitions in the original automata, trying to find such runs where all the original automata use the same transition symbols. Therefore, one can do the same run with the transition symbols for all of the finite automata at once in the product—assuming there is such a run—combining the original states from the individual runs to tuples called product states and adding them to the generated state space. For two automata, the classical algorithm starts with an initial state of both automata, make pairs from them and try to find transitions leading from these initial states in the original automata with the same symbol. Generate new pairs for found transitions, find their transitions, and so on. Every

product state represents an intersection of languages of two corresponding states in the original automata.

Unfortunately, the synchronous product construction is expensive on computational time as one needs to generate a vast number of product states during the process. For two automata, the state space can increase quadratically, for more automata, the product state space increases exponentially according to the number of used automata and the number of their states. Furthermore, often there are large parts of the generated state space which cannot accept any words (non-terminating states), yet are still generated because of the corresponding transition symbols.

We will try to optimize this process by reducing the number of generated product states and their transitions for both generating the synchronous product and deciding its emptiness test. We focus mainly on decision-making about the satisfiability problem-solving the emptiness of an intersection of two (eventually even multiple) finite automata. We consider length abstraction over the initial finite automata while trying to predict which product states cannot lead to any accepting states.

When the lengths of words recognized by the languages of the current states are not compatible with each other—the original languages of the corresponding states cannot accept a word of the same lengths—there is definitely no transition from this product state leading to accepting the same word in both original automata. We can omit such states from our generated product. Consequently, this removes the need to even consider their potential successor states, which are generated normally. By doing this, we trim the generated product state to only states whose corresponding original states languages can accept words of the same lengths. Even though there might still be states which do not lead to any accepting state in the final product, this simple optimization already trims a substantial parts of the normally generated synchronous product state space reasonably often.

Computing length abstraction over the languages of finite automata (and over individual states in the automata in particular) is accomplished using lasso automata (handle and loop automata)—deterministic finite automata with a unary alphabet (similar as in [1]). They consist of a *handle* (a sequence of states from the initial state) and a single *loop* (resolving the cycles in the original automaton) with a few accepting states along the way, hence the name of a handle and loop or lasso automaton. Create one by taking the initial automaton, consider all transition symbols as a single transition symbol and determinize this automaton.

What you get is an automaton accepting every length of any word recognized by the language of the initial automaton. Consequently, it is easy to compute semi-linear set (formulae¹) for the allowed lengths of words, which can be effectively processed and compared using SMT solvers. We are computing these formulae for individual product states (precisely for the corresponding states in the original automata), checking their satisfiability and consequently construct only those product states for which the length test resolves as satisfiable.

The second optimization approach we consider is the computation of Parikh images for potential product states. Computing Parikh image as a set of constraints for given product state to be included in the generated product. With Parikh images of such potential product states, we have additional information about the current state under consideration and can more precisely determine whether to include such state in the product state space. However, the Parikh image computation in itself is a demanding operation and the eval-

¹disjunction of linear equations

uation of satisfiability of mentioned constraints for every product state requires additional computation time.

In this case, it is necessary to decide whether the trade-off of unoptimized basic algorithm generating larger product state space requiring less computation time for reduced product state space generated by our optimized algorithms using Parikh images with additional computation time requirements is worth our attention. We say it is. For certain operations over the automata, the product state space size is crucial, considering we may need to work with the same product multiple times or simply need to execute a single operation on the product, reduced state space can spare extensive amounts of computation time further down the processing line. Furthermore, generating smaller state space using our Parikh image optimization can improve computation time for the sole product generation algorithm in case substantial parts of otherwise generated state space are pruned due to unsatisfiable constraints defined by Parikh image for any given product state in the process of generating the final product or even when the whole product is proved to be empty, which can be quickly determined by our optimization, whereas the classic unoptimized algorithms would proceed to generate useless fragments of suppositional product.

Our another optimization uses mintermization as a different approach to abstracting state languages of initial automata. We compute minterms, which can be used instead of transition symbols while retaining all information about the automata to compute Parikh images and other optimization abstractions in less computation time.

We have implemented these optimizations and experimented with several different automata, tried various combinations of them, generated their products and tried to solve their emptiness test, focusing mainly on the number of trimmed product states in the process. For certain types of automata of certain qualities, this optimization process works really well.

Our proposed algorithm cannot remove any product states leading to an actual accepting state. Therefore, it is impossible to accidentally trim product states leading to some accepting states and change the intersection language in the end. The language we get is the same as the one generated from the naive product construction algorithm. Consequently, is it completely secure to use our optimization with any kind of automata for any kind of uses.

The contribution of this work can be summarized as follows:

1. heuristics trimming the generated state space of finite automata operations based on length abstraction, Parikh image computation, mintermization, and
2. implementation and experimental evaluation of said heuristics.

Chapter 2

Preliminaries

Let us clarify a few definitions and terms often used throughout this paper. The following definitions are mostly adapted from [7] or [19].

Alphabet is a finite, non-empty set denoted by Σ . Elements of an alphabet are called *symbols* or *letters*. A finite, possibly empty sequence of symbols over an alphabet is a *word* w from the set of all words Σ^* over an alphabet Σ .

Definition 2.0.1 (Deterministic finite automaton)

A *deterministic finite automaton (DFA)* is a 5-tuple $A = (Q, \Sigma, \delta, I, F)$, where:

- Q is a **non-empty set of states**,
- Σ is an **input alphabet**,
- δ is a **transition function**: $Q \times \Sigma \rightarrow Q$,
- $I \in Q$ is the **initial state**, and
- $F \subseteq Q$ is a **set of final states**.

A *run* of A on input $a_0a_1a_2\dots a_{n-1}$ is a sequence $q_0 \xrightarrow{a_0} q_1 \xrightarrow{a_1} q_2 \xrightarrow{a_3} \dots \xrightarrow{a_{n-1}} q_n$, such that $q_i \in Q$ for $0 \leq i \leq n$ and $\delta(q_i, a_i) = q_{i+1}$ for $0 \leq i \leq n-1$. A run is *accepting* if $q_n \in F$. The automaton A accepts a word $w \in \Sigma^*$ if it has an accepting run on input w . A *language* recognized by finite automaton A is a set $L(A) = \{w \in \Sigma^* | w \text{ is accepted by } A\}$. A single transition from transition function δ is denoted as $q \xrightarrow{a} q'$ if $q' \in \delta(q, a)$ and means *one can get from state q to state q' with a transition symbol a* . For every state, DFA has at most one transition for a given symbol.

Definition 2.0.2 (Non-deterministic finite automaton) A *non-deterministic finite automaton (NFA)* is a 5-tuple $A = (Q, \Sigma, \delta, I, F)$, where Q , Σ and F are as for DFA and:

- δ is a **transition relation**: $\delta : Q \times \Sigma_\epsilon \rightarrow P(Q)$, where $\Sigma_\epsilon = \Sigma \cup \epsilon$ and $P(Q) = \{R | R \subseteq Q\}$ is a set of subsets of Q , and
- $I = \{q | q \in Q\}$ is a **non-empty set of initial states**.

Consequently, DFA has exactly one run on a given word from initial state to one of the accepting states (or nonterminating states in case the word is not accepted by the

automaton at all). For every state and its transition symbol $P(Q) \in \delta(q, a)$ is a singleton. For example, $\delta(q_1, a) = \{q_1, q_2\}$.

Two finite automata A and B are said to be *equivalent* when both accept the same language: $L(A) = L(B)$.

For every NFA A exists a corresponding DFA B . *Determinization* is a process of converting such NFA to DFA.

Definition 2.0.3 (Powerset (subset) construction) *The powerset construction is a method for creating a corresponding deterministic finite automaton from its equivalent non-deterministic finite automaton. Produces finite automaton A' , where $Q' = 2^Q$, $F' = \{S \in Q' \mid S \cap F \neq \emptyset\}$, $I' = I$ and for $S \in Q' : \delta'(S, a) = \bigcup_{s \in S} \delta(s, a)$.*

Definition 2.0.4 (Product construction) Operations *on automata A_1 and A_2 yield a result – a product A as a 5-tuple deterministic finite automaton $A = (Q, \Sigma, \delta, I, F)$.*

Given two NFAs $A_1 = (Q_1, \Sigma, \delta_1, I_1, F_1)$ and $A_2 = (Q_2, \Sigma, \delta_2, I_2, F_2)$ over the same alphabet Σ , we can define:

- *a set of states $Q = Q_1 \times Q_2$,*
- *a transition relation $\delta : Q \times \Sigma \rightarrow P(Q)$,*
- *a set of initial states $I = I_1 \times I_2$, and*
- *a set of accepting states $F = F_1 \times F_2$.*

The transition relation is described as $\delta = ([q_1, q_2], a) = \delta_1(q_1, a) \times \delta_2(q_2, a)$. For pairs of states q_1 and q_2 from A_1 and A_2 , respectively, and a common transition symbol a of transitions $q'_1 \in \delta_1(q_1, a)$ and $q'_2 \in \delta_2(q_2, a)$, we denote a single product transition as $[q_1, q_2] \xrightarrow{a} [q'_1, q'_2]$, where $[q'_1, q'_2] \in \delta([q_1, q_2], a)$ for the corresponding states $[q_1, q_2]$ and $[q'_1, q'_2]$ in A are called product states.

Focusing mainly on *intersection* of automata, the product construction tells that

$$L(A) = L(A_1) \cap L(A_2).$$

To solve the *emptiness test*, we test

$$L(A) = \emptyset.$$

Input : NFA $A_1 = (Q_1, \Sigma, \delta_1, I_1, F_1)$,
 NFA $A_2 = (Q_2, \Sigma, \delta_2, I_2, F_2)$
Output: NFA $(A_1 \cap A_2) = (Q, \Sigma, \delta, I, F)$ with $L(A_1 \cap A_2) = L(A_1) \cap L(A_2)$

```

1  $Q, \delta, F \leftarrow \emptyset$ 
2  $I \leftarrow I_1 \times I_2$ 
3  $W \leftarrow I$ 
4 while  $W \neq \emptyset$  do
5   pick  $[q_1, q_2]$  from  $W$ 
6   add  $[q_1, q_2]$  to  $Q$ 
7   if  $q_1 \in F_1$  and  $q_2 \in F_2$  then
8      $\lfloor$  add  $[q_1, q_2]$  to  $F$ 
9   forall  $a \in \Sigma$  do
10    forall  $q'_1 \in \delta_1(q_1, a), q'_2 \in \delta_2(q_2, a)$  do
11      if  $[q'_1, q'_2] \notin Q$  then
12         $\lfloor$  add  $[q'_1, q'_2]$  to  $W$ 
13      add  $[q'_1, q'_2]$  to  $\delta([q_1, q_2], a)$ 

```

Algorithm 1: Classic product construction

Chapter 3

State Language Optimizations

In this chapter, we will introduce several optimizations of state language. We aim to optimize operations on finite automata such as product construction, complement computation, minimalization or determinization and inclusion test. Furthermore, we want to introduce state language optimizations which work for different automata structures. E.g., operations on transducers, operations with alternating automata such as its emptiness or conversion of alternating automaton to its NFA representation, conversion of finite automata to flat automata, etc.

Not only the classic operations are generally useful, but both transducers and alternating automata are operations often used, beside others, for example, in verification. Our optimization methods therefore have the ability to improve substantially important processes used throughout multitude of places and fields of study as well as in praxis.

We perform experiments with our optimization methods mostly on product construction of two NFAs and our optimization method algorithms are consequently introduced on these algorithms, improving the naive product construction algorithm to generate optimized products. Nevertheless, our optimization techniques are to be used in various fields of automata theory and allow us to optimize other, more complex problems. We chose product construction as our benchmarking operation on automata for its straightforwardness allowing us to see clearly pruning capabilities of our proposed optimization methods, even though the naive product construction algorithm complexity can be at most *only* quadratic and is not that expensive in terms of computational time. Our optimization methods are applicable for more complex operations for exponential constructions (determinization, minimization, emptiness of alternating automaton, ...) and worse.

3.1 State Language Optimization with Length Abstraction

Our task is to try to minimize the number of generated states when trying to resolve the product construction of automata intersection and test its emptiness. One possible solution is looking for lengths of words accepted by both automata—testing whether both automata recognize words of the same lengths. Afterwards, we check the original transition symbols for generating new product states¹. Consequently, we can resolve the emptiness test of some intersections very quickly and optionally optimize the product construction, when we need to generate the whole product.

¹So we do not get non-empty intersection results when there is no word both original automata actually accept and only their lengths correspond.

We will explain our chosen approach to the problem of optimizing product construction and deciding its emptiness test using length abstraction, but first some rudimentary knowledge on length abstraction is needed.

3.1.1 Length Abstraction Represented by Lasso Automata

Our chosen approach to the problem of optimizing product construction and deciding its emptiness test includes using length abstraction over the finite automata to try to guess which product states do not lead to any final states and consequently can be omitted, and the following states do not need to be generated at all.

Length abstraction generalizes the language recognized by the initial automaton by considering only the possible lengths of words accepted by the automaton. It is an over-approximation of the language accepted by the original automaton. For us, this means if a word is not accepted by the length abstraction automaton, it cannot be accepted by the initial automaton either.

The length abstraction automaton is represented by a so-called lasso automaton. Let us demonstrate creation of the lasso automaton on the following simple non-deterministic finite automaton A_1 , which we will continue to use in this paper to depict our optimization algorithm.

$$A_1 = (\{q_0, q_1, q_2, q_3, q_4, q_5\}, \{0, 1\}, \delta_1, \{q_0\}, \{q_4\})$$

Transition relation δ_1 is depicted in Figure 3.1.




Figure 3.1: Non-deterministic finite automaton A_1

NFA A_1 is a non-deterministic finite automaton (see state q_1) and accepts more than one input symbol. Due to the fact that we work only with recognized word lengths, we can substitute automaton alphabet with unary alphabet of single input symbol (we have chosen $*$ for demonstration purposes)². Then, we can compute lasso automaton for our original automaton A_1 with unary alphabet, which is its deterministic equivalent.

²Even though we do not actually need any particular input symbol, we use $*$ here as an example to depict the process. In general, all we need to know is that there is a transition between two states, the transition symbols are not significant for our optimization algorithm.

$$\Sigma = \{0, 1\} \longrightarrow \Sigma' = \{*\}$$




Figure 3.2: Non-deterministic finite automaton A_1 with unified transition symbols

We start the determinization process on our updated automaton. For the final lasso automaton A'_1 for the original automaton A_1 , see Figure 3.3. This automaton now accepts any words of lengths of words recognized by the original automaton. We will use these lengths in the process of constructing the product.³




Figure 3.3: Lasso automaton A'_1 for the original NFA A_1

³You can notice this lasso automaton looks different to what is depicted in Section 3.1.1. It is caused by our optimization with generating only one single lasso automaton per original automaton. The generated automaton is valid and the fact there are actually two separate automata with one even being inaccessible does not raise any issue for us. The reason for this behaviour will be further explained in Section 3.1.2.

3.1.2 Single Lasso Automaton for Each Original Automaton

When we are constructing a product, we do not want to regenerate the lasso automaton L for each new product state q . This is inefficient. Therefore, our algorithm generates L only once—state by state—checking every time, whether the new state l is not already present in a set of states Q_L of L .

Due to the nature of lasso automata, the successive product states generate the lasso automata very similar to L . We just need to append new states to Q_L . As a result, we will work with only two lasso automata (possibly with multiple loops and/or multiple handles)—one for both automata whose intersection is computed.

If $l \notin Q_L$, we add l to Q_L and continue with the following states l' until we either create an entirely new loop in L or generate $l' \in Q_L$. If $l \in Q_L$, we can stop generating l' for q as $\forall l' (l' \in Q_L)$.

3.2 Product Construction with Length Abstraction Optimization

The core of the product construction algorithm remains unchanged, but there are a few differences. The Algorithm 2 shows how we alternate the original product construction algorithm to optimize it and resolve emptiness test for each *branch* of the potential product automaton.

Input : NFA $A_1 = (Q_1, \Sigma, \delta_1, I_1, F_1)$,

NFA $A_2 = (Q_2, \Sigma, \delta_2, I_2, F_2)$

Output: NFA $(A_1 \cap A_2) = (Q, \Sigma, \delta, I, F)$ with $L(A_1 \cap A_2) = L(A_1) \cap L(A_2)$

```

1  $Q, \delta, F \leftarrow \emptyset$ 
2  $I \leftarrow I_1 \times I_2$ 
3  $W \leftarrow I$ 
4  $sat \leftarrow False$ 
5  $solved \leftarrow \emptyset$ 
6 while  $W \neq \emptyset$  do
7   picklast  $[q_1, q_2]$  from  $W$ 
8   add  $[q_1, q_2]$  to  $solved$ 
9    $sat \leftarrow \text{satisfiable}([q_1, q_2])$ 
10  if  $sat$  then
11    add  $[q_1, q_2]$  to  $Q$ 
12    if  $q_1 \in F_1$  and  $q_2 \in F_2$  then
13      add  $[q_1, q_2]$  to  $F$ 
14    forall  $a \in \Sigma$  do
15      forall  $q'_1 \in \delta_1(q_1, a), q'_2 \in \delta_2(q_2, a)$  do
16        if  $[q'_1, q'_2] \notin solved$  and  $[q'_1, q'_2] \notin W$  then
17          add  $[q'_1, q'_2]$  to  $W$ 
18          add  $[q'_1, q'_2]$  to  $\delta([q_1, q_2], a)$ 

```

Algorithm 2: Product construction with length abstraction

We will call W from line 3 a work set. It stores the potential product states prepared for testing for satisfiability and other processing, which we pick from W one by one⁴.

The optimization process starts when a product state q is picked from W . Instead of immediately generating new successive product states, we test q for satisfiability of length constraints of recognized words from q . On line 9, we check for satisfiability of formulae generated from q . If the process shows formulae are satisfiable, i.e., there will be an accepting run using q (see line 10), we add q to Q , possibly to F and generate its successive product states q' .

```

1 Function isLengthAbstractionSatisfiable(formulaeForFA1, formulaeForFA2):
   Data: Input length formulae of potential product state we are solving the
           satisfiability for (for all possible accept states combinations).
   formulaeForFA1: Formulae for the first finite automaton,
   formulaeForFA2: Formulae for the second finite automaton.
   Result: bool: True if satisfiable, False otherwise.
2   smtAdd(VariableForFormulaForFA1  $\geq$  0, VariableForFormulaForFA2  $\geq$  0)
3   for formulaForFA1  $\in$  formulaeForFA1 do
4       for formulaForFA2  $\in$  formulaeForFA2 do
5           smtPush()
6           smtAdd(formulaForFA1.handle + formulaForFA1.lasso *
                   VariableForFormulaForFA1 = formulaForFA2.handle +
                   formulaForFA2.lasso * VariableForFormulaForFA2)
7           sat  $\leftarrow$  smtCheck()
8           if sat or sat = unknown then
9               return True
10          smtPop()
11 return False

```

Algorithm 3: Check satisfiability using length abstraction algorithm with SMT solver

The formulae are generated using lasso automata for both original automata. For every state, we get one or more formulae in the form $\varphi : \exists k (|w| = a + b \cdot k)$, where $|w|$ is a length of a recognized word, a is the length of a handle to a certain accepting state, and b is the length of a loop to return to this particular accepting state going through the loop. k is the number of cycles through the loop states until a word ends in an accepting state. When multiple depicted formulae are present (because there are more accepting states in the lasso automaton), we append these formulae with *logical or* (\vee), then compare these with the formulae from the other lasso automaton for the other initial finite automaton using SMT solver.

To better demonstrate our solution, the second automaton we will be working with is a finite automaton A_2 from Figure 3.4.

⁴In spite of the fact that more approaches are valid, we strongly recommend picking the last added product state from the work set W (see line 7)—using Depth-first Search for a graph algorithm—as this allows us to quickly advance through the automaton and get to any final state faster—in case we just want to know whether automata have a non-empty intersection, this change will get us the answer most of the time in less steps. It works even better with implemented satisfiable state skipping optimization, explained in Section 3.2.1.

$$A_2 = (\{s_0, s_1, s_2, s_3\}, \{0, 1\}, \delta_2, \{s_0\}, \{s_3\})$$

Transition relation δ_2 is depicted in Figure 3.4.




Figure 3.4: Non-deterministic finite automaton A_2

In Figure 3.5, there is its lasso automaton A'_2 , which we will be using together with the lasso automaton shown earlier in Figure 3.3 for computation of recognized word lengths.




Figure 3.5: Lasso automaton A'_2 for the original NFA A_2

For automaton A_1 for the initial state (we start computing lengths as if the state q_0 is the initial state) from A'_1 , we get quantifier-free formula φ^5 . For automaton A_2 for the

⁵This formula consists of two independent formulae describing there are more possible lengths for accepted words from the same initial state (leading to two independent accepting states in the automaton).

initial state from the lasso automaton A'_2 , we get formula ψ^6 .

$$\varphi : \exists k(|w| = 2 \vee |w| = 4 + 2 \cdot k)$$

$$\psi : \exists l(|w| = 2 + 1 \cdot l)$$

When we compare φ and ψ , we get:

$$\exists k \exists l (2 \vee 4 + 2 \cdot k = 2 + 1 \cdot l)$$

We try to find values of k and l such that some of the expressions on the left and on the right side of the equation are equal. We pass this equation to SMT solver to solve its satisfiability. Returns *sat* when satisfiable (*sat* is set to *True*) and *unsat* when unsatisfiable (*sat* is set to *False*). If *unsat* is returned, we can stop generating this *branch* of a NFA as we know for sure there cannot be a word which is accepted by both of these automata, when there is even no word fulfilling the length requirements. In this case, we have successfully reduced the generated state space by omitting this particular tested product state and any further product states, which would be later normally generated from these product states and its successors (assuming the transition symbols correspond).




Figure 3.6: Constructed product automaton with depiction of our optimization.

In Figure 3.6, we can see the product of A_1 and A_2 being constructed using our optimization. Red states represent tested states that are resolved as unsatisfiable for computed length formulae and therefore the algorithm omits any successive product states—dashed states (such as q_4s_2 or q_3s_2), which are generated in the basic naive product construction algorithm. The green state is satisfiable and also represents accepting states in both automata. Here, we have found a possible solution accepted by both original automata. If we desire to resolve only the product emptiness test, we can stop the execution of the algorithm here as we have found one accepting state—automata have non-empty intersection.

⁶We are using variable l here instead of k to emphasize variables from different formulae are not dependent on each other—they correspond to various accepting states.




Figure 3.7: Final product generated by our product construction algorithm optimizing its process by omitting unnecessary states on the way.

As you can notice in Figure 3.7, the product generated by our algorithm has only 4 product states in comparison to 9 product states generated by the classic algorithm.

When we get the result of formulae being unsatisfiable, we do not need to construct any following states and check their products for satisfiability and therefore, we are able to determine whether such *branches* of automata have an empty intersection and we do not need to consider them in the product construction. The emptiness test is successfully accomplished, and we determined that for this *branch* there cannot exist any word accepted by both automata and consequently by their intersection. When we get unsatisfiable result for every branch of the automaton (i.e., no *branch* can lead to any accepting state) even if by inspecting transition symbols it looks like there could be a non-empty intersection, we can say that such input automata have an empty intersection and product construction will be very quick in that case—this is where our optimization dominates.

A note of caution. It is important to understand that we are working only with possible word lengths and therefore when we test the emptiness of intersection of automata, we can resolve only such intersections that words lengths are not accepted by both automata. When the test shows there could be some words of certain length accepted by both automata and for that reason by their intersection too (the result of the length abstraction satisfiability check is *sat*), we cannot be sure there truly are any words accepted by both automata with their intersection non-empty, because there may be words of the suggested length, but it may be a different word for each automaton (which differ from one another in the containing symbols or their position in the word). For resolving such cases, we have to proceed with the classic algorithm steps to produce product states according to their original transition symbols, not only by comparing the possible words lengths. With certainty, we can omit only the cases where the length abstraction satisfiability check returns *unsat*.

3.2.1 Optimization with Skipping Satisfiable States

When we take a new product state q from work set W and check for satisfiability with formulae for q being satisfiable, it is time to add to W all of the possible successive product states q' . When q generates only a single q' , we can say with certainty formulae for q'

are satisfiable too as there is only a single branch in the automaton leading from q to an accepting state (through q'). Product state q' is *skippable*, if exists a satisfiable q whose only successor is q' . We add q' to W with the information of being skippable. If q' is already in W , we append the information to q' in W .

We skip checking for the satisfiability when we pick q' from W . We can immediately check for final states and generate the successive product states. This optimization saves us generating the length abstraction formulae for q' and testing the formulae in the SMT solver for their satisfiability and even possibly reducing the number of states generated for both our lasso automata.

If our original automata have long lines (with non-splitting branches), this will prove extremely useful, because only a few proper iterations with formulae computing and executing SMT solver will be executed. In Algorithm 4 is depicted the application of skipping satisfiable states. The line 9 from Algorithm 2 is substituted with the contents of Algorithm 4.

```

1 if not skippable( $[q_1, q_2]$ ) then
2   |  $sat \leftarrow \text{satisfiable}([q_1, q_2])$ 
3 else
4   |  $sat \leftarrow \text{True}$ 

```

Algorithm 4: Substitution of line 9 in Algorithm 2 with skipping satisfiable states.

The only change is a test for every checked product state q , which decides whether q can be skipped, if it cannot give us any information which we do not have yet. You can see that we proceed with SMT solver satisfiability check only for q which are generated from the product states with multiple transitions generating q and at least one more product state (in general at least two new potential product states). If only one q is generated, we skip the satisfiability check for q and continue to generating its successive states immediately.




Figure 3.8: Constructed product automaton with depiction of skipping states optimization

You can notice there is one skippable state in the former example, which had to be evaluated and tested for satisfiability earlier. The blue state in Figure 3.8 is such a skippable state. In our case for state q_5s_2 , when only one new state is generated from state q_4s_3 while this state is resolved as satisfiable, newly generated product state has to be satisfiable as well, because the check for q_4s_3 already considered the state q_5s_2 as its only way to any accepting state resolving in *sat* check for q_4s_3 actually.

When we have a series of such states, though, we can highly optimize generating the whole branch with only one initial check for satisfiability. In real world examples, there are often automata with long branches splitting into multiple branches only occasionally. We will check for satisfiability for all of the initial states of each new branch and then either omit the entire branch (if *unsat* is returned) or skip checking satisfiability in the entire branch (if *sat* is returned).

3.2.2 Resolving Length Abstraction Satisfiability without SMT Solver

```

1 Function isLengthAbstractionSatisfiable(formulaeForFA1, formulaeForFA2):
   Data: Input length formulae of potential product state we are solving the satisfiability
           for (for all possible accept states combinations).
   formulaeForFA1: Formulae for the first finite automaton,
   formulaeForFA2: Formulae for the second finite automaton.
   Result: bool: True if satisfiable, False otherwise.
2 for formulaForFA1  $\in$  formulaeForFA1 do
3   for formulaForFA2  $\in$  formulaeForFA2 do
4     if formulaForFA1.handle = formulaForFA2.handle then
5       return True
6     else if formulaForFA1.handle > formulaForFA2.handle then
7       sat  $\leftarrow$  solveForOneHandleLonger(formulaForFA1, formulaForFA2)
8       if sat then
9         return True
10    else
11      sat  $\leftarrow$  solveForOneHandleLonger(formulaForFA2, formulaForFA1)
12      if sat then
13        return True
14 return False

```

Algorithm 5: Check satisfiability using length abstraction algorithm without SMT solver

3.3 Product Construction Optimization with Parikh Image Computation

This section presents a product construction optimization using Parikh images of finite automata. Parikh image optimization provides more information about the finite automata the product is computed for than simple length abstraction, which allows us to more precisely determine whether the generated product has non-empty intersection. However, the

```

1 Function solveForOneHandleLonger(formulaForFA1, formulaForFA2):
   Data: Input length formulae of potential product state we are solving the satisfiability for
           (for concrete accept states combination).
   formulaForFA1: Formula for the first finite automaton,
   formulaForFA2: Formula for the second finite automaton.
   Result: bool: True if satisfiable, False otherwise.
2   formulaForFA1.handle  $\leftarrow$  formulaForFA1.handle - formulaForFA2.handle
3   formulaForFA2.handle  $\leftarrow$  0
4   if formulaForFA1.lasso = 0 and formulaForFA2.lasso = 0 then
5     | return False
6   else if formulaForFA2.handle = 0 then
7     | return False
8   else if formulaForFA1.lasso = 0 then
9     | currentIteration  $\leftarrow$  0
10    | while currentIteration  $\leq$  formulaForFA1.handle do
11      | if currentIteration = formulaForFA1.handle then
12        | return True
13      | else
14        | currentIteration  $\leftarrow$  currentIteration + formulaForFA2.lasso
15      | return False
16  else
17    | gcd  $\leftarrow$  getGCD(formulaForFA1.lasso, formulaForFA2.lasso)
18    | if gcd = 1 then
19      | return True
20    | else
21      | y  $\leftarrow$  -formulaForFA1.handle
22      | while y < gcd do
23        | y  $\leftarrow$  y + formulaForFA1.lasso
24      | if y mod gcd = 0 then
25        | return True
26      | else
27        | return False
28  return False

```

Algorithm 6: Solve satisfiability of length abstraction formulae for one handle longer.

Parikh image computation itself consumes a considerable amount of computational time for some of the more extensive finite automata. The question is, whether the added computation time compensates for more precise product generation with higher product states pruning capabilities.

In this section, we will introduce an algorithm for Parikh image computation used for computing Parikh image of each potential product state to decide the satisfiability of Parikh image constraints for this particular product state. If the constraints are proved to be satisfiable, the considered product state is added to the generated product, otherwise, the potential product state is omitted and no additional potential product states accessible from the omitted product state in one step are added to the queue for testing for their satisfiability.

3.3.1 Parikh Image

We derive our Parikh image construction from the Parikh's theorem [15] described in [8], creating a semilinear Parikh image formulae for the given regular language as a set of Parikh images for each word in the language. However, our usage of Parikh image of some regular language (and therefore of the corresponding finite automaton recognizing such regular language) is restricted to determine the satisfiability of Parikh image formulae. That is, Parikh image constrains which we use in SMT solver to represent the abstraction of the current potential product state language.

Our Parikh image formulae consist of the following constraints (clauses), in conjunctive normal form. For each potential product state, there exists exactly one our formula of Parikh image describing its regular language. We ask the SMT solver whether the Parikh image constraints for corresponding states in the original automata (one state per automaton) are compatible with each other. This ensures that we construct only those product states which satisfy the Parikh image constraints, otherwise we deem such potential product states redundant and such states can be pruned.

Given an NFA $A = (Q, \Sigma, \Delta, q_0, F)$, Parikh image constraints representing clauses in the classic Parikh image formula are constructed as follows:

1. Foremost, we define a variable u_q for each transition $t \in \Delta$. We assign values to u_q as follows:
 - $u_q = 1$ for the initial state q_0 ,
 - $u_q \in \{0, -1\}$ for each state $q \in F$ and
 - $u_q = 0$ for every other state $q \in Q \setminus (\{q_0\} \cup F)$.
2. Second, we define a variable y_t for each transition $t \in \Delta$ such that $y_t \geq 0$.
3. We can now present an equation introducing a connection between u_q and y_t for each automaton state $q \in Q$ as follows:

$$u_q + \sum_{t \in \Delta_q^+} y_t - \sum_{t \in \Delta_q^-} y_t = 0.$$

where Δ_q^+ is a set of ingoing transitions $\Delta_q^+ = \{(q', a, q) \in \Delta\}$ and Δ_q^- is a set of outgoing transitions $\Delta_q^- = \{(q, a, q') \in \Delta\}$ from the given state q .

4. Furthermore, we declare a variable $\#_a$ for each transition symbol $a \in \Sigma$ such that $\#_a = \sum_{t=(q,a,q') \in \Delta} y_t$ to express the number of uses of a given symbol a is consistent with the number of used y_t with a .
5. Last, but not least, we make sure the regular language expressed by Parikh image preserves the connectedness of the finite automaton. For this reason, yet another variable z_q for each state $q \in Q$ is introduced. z_q represents the length of the path from q_0 to q in a spanning tree of the subgraph with $y_t \geq 0$.

If the state q is an initial state, we add a constraint $z_q = 1 \wedge y_t \geq 0$. Otherwise,

$$(z_q = 0 \wedge \bigwedge_{t \in \Delta_q^+} y_t = 0) \vee \bigvee_{t \in \Delta_q^+} (y_t \geq 0 \wedge z_{q'} \geq 0 \wedge z_q = z_{q'} + 1).$$

We gain a qualifier-free formula φ in Presburger arithmetic:

$$\exists u_{q_1}, \dots, u_{q_n}, z_{q_1}, \dots, z_{q_n}, y_{t_1}, \dots, y_{t_m} : \varphi$$

where $n = |Q|$ is the number of states and $m = |\Delta|$ is the number of transitions in the finite automaton.

Reduced Parikh Image

The presented Parikh image works as intended. Nevertheless, the described Parikh image computation requires extensive resources and computation time. However, we use Parikh image only for determining the emptiness of the product. Given that most of the computation time takes evaluation of satisfiability of Parikh image formula φ , we try to minimize the number of Parikh image formula clauses SMT solver needs to evaluate in order to determine satisfiability of φ .

Consequently, we infer our reduced Parikh image from the above shown Parikh image to further optimize Parikh image computation. We strip Parikh image of for our purposes unnecessary constraints and unifying initial states as well as accept states for initial finite automata.

Our reduced Parikh image consists of the following clauses:

1. We use the clause 1, except now we restrict u_q for each final state to have only the value -1 , i.e.:

$$u_q = -1 \text{ for each state } q \in F.$$

We can perform this reduction, because we know for sure that by unifying final states of the automaton into one abstract final state, there will be exactly only one final state where all words accepted by the automaton end, but none passes through this state earlier.

2. The clause 2 and 3 remain unchanged, the same holds for clause 4.
3. However, we completely omit the clause for z_q which ensures the connectedness of the Parikh image representation of finite automaton. The reason is that, as we have found out, the difference in pruning capabilities of Parikh image with or without the clause 5 on our benchmark automata is insignificant in comparison to the computation time spared by removing this clause (see Chapter 4).

The reason the clause 5 is so demanding computation time-wise is that the whole clause have to be always recomputed for each single state Parikh image is computed for. Furthermore, the clause itself is complex enough for even simple automata and SMT solvers need extensive resources to compute Parikh images with this clause in consideration.

Even then, if we require ensuring that the reduced Parikh image represents the connectedness of the finite automaton, we can include this clause, but, thanks to our unification of initial and accept states, we change it as follows to reflect our initial and accept state unification changes:

The constraint for when q is an initial state ($z_q = 1 \wedge y_t \geq 0$) remains unchanged. However, for every other state, we remove the possibility of $y_t = 0$ and $z_{q'} = 0$ in the second half of the clause. The clause looks like this:

$$(z_q = 0 \wedge \bigwedge_{t \in \Delta_q^+} y_t = 0) \vee \bigvee_{t \in \Delta_q^+} (y_t > 0 \wedge z_{q'} > 0 \wedge z_q = z_{q'} + 1).$$

Our goal is to reduce the number of constraints the SMT solver needs to compute for each potential product-state. We focus on several optimizations such an incremental SMT solving,

Due to how we have reduced our Parikh image used for automata state language abstraction, we work only with finite automata with a single initial state and a single accept state. However, we can easily convert any finite automaton into the required format with adding two new states to each input automaton. One for a new initial state from which one can transition to all previous initial states and one for a new accept state to which leads all previous accept states. The previous initial and accept states are changed to common automata states.

Satisfiability of Multiple Parikh Image Formulae Simultaneously

So far, we have shown how to compute Parikh image for a single finite automaton to represent said automaton with a single formula. We want to use this formula in such a way that would allow us to decide satisfiability of those formulae for multiple automata simultaneously when the formulae are combined into a single formula which we can decide its satisfiability for.

Given at least two finite automata $A_i \in A$, we can compute Parikh image formulae $\varphi_i \in \Phi$ for each finite automaton where $1 \leq i \leq n$ for $n = |A|$ is the number of finite automata. Each of these formulae $\varphi_i \in \Phi$ represents exactly one of the finite automata $A_i \in A$. Therefore, each of the formulae φ_i by itself is satisfiable in some interpretation I_i such that $I_i \models \varphi_i$ for the automaton A_i where I_i evaluates variables in φ_i such that the formula describes words accepted by the automaton A_i ⁷.

We use this features of Parikh image to determine satisfiability of multiple finite automata as follows. If each formula φ_i is satisfiable, we want to know whether a combination of multiple formulae $\varphi_i, \varphi'_i, \dots$ is satisfiable at the same time for the same interpretation I . However, to maintain the languages of specific automata distinguishable, we label each

⁷Our Parikh image is an overapproximation of the accepted language of a given finite automaton. Therefore, there could exist such interpretation $I \models \varphi_i$ which describe words not accepted by the original finite automaton A_i . It is a trade-off of precise representation of given automaton for faster computation of Parikh image.

variable u_q, y_t (optionally, z_q , too) for each φ_i according to i : $u_{iq}, y_{it} (z_{iq})$. The only exception are variables $\#_a$ which in contrary are bound to transition symbols $a \in \Sigma$ common to all finite automata $A_i \in A$.

Now, we find interpretation I such that:

$$I \models \bigwedge_{1 \leq i \leq |A|} \varphi_i$$

which would mean there are words accepted by all the Parikh image formulae $\varphi_i \in \Phi$ simultaneously and therefore by all the automata from A . Consequently, the automata product would be non-empty.

3.3.2 Optimization Algorithm Using Parikh Images

In this section, we introduce the basic algorithm using Parikh image computation to construct the product of the intersection of finite automata. The algorithm resembles length optimization algorithm from Algorithm 2. However, we compute Parikh image formulae and determine their satisfiability instead of generating lasso automata and determining satisfiability of length abstraction formulae now to optimize product construction.

We use Parikh image formulae to determine whether potential product states are to be added to the generated product (in case the Parikh image formulae for each automaton $A_i \in A$ are satisfiable simultaneously) or omitted (in case any Parikh image formula φ_i for a finite automaton $A_i \in A$ is unsatisfiable simultaneously with the Parikh image formulae φ_j where $j \in \{1 \leq j \leq |A|\} \setminus \{i\}$ for the remaining finite automata $A_j \in A$).

We can see our proposed algorithm using Parikh image computation to optimize product construction in the Algorithm 7. Similarly to the length abstraction algorithm, we start with the initial states (our abstract initial state, as described in Section 3.3.1) of all finite automata $A_i \in A$, compute Parikh image formulae for each A_i combined into a single formula φ . If φ is unsatisfiable (each Parikh image formulae simultaneously), the product is empty and we can stop the product generation at once. Otherwise, the formula is satisfiable and the corresponding product state is added to the generated product P . We proceed to generate the consecutive potential product states. We set the initial states for Parikh image formulae computation to the current state for each automaton A_i for each potential product state and recompute the combined Parikh image formula. We iterate over potential product states from W (see line 6).

The function `areParikhImageFormulaeSatisfiable` computes Parikh image formulae, determines their satisfiability and returns the result as a boolean value. We are only interested in the satisfiability test result because we do not need any additional information from the computed formulae. Therefore, a simple boolean value is sufficient. The result of the satisfiability test is used further in the algorithm to determine whether the product state is added to the generated product and append consecutive potential product states to W . The Parikh image is computed as it is explained in Section 3.3.1.

Optimization with Skippable States

Same as for the length abstraction algorithm from Algorithm 2, we can make use of skipping satisfiable product states optimization. When Parikh image is evaluated as satisfiable for some potential product state q and such state generates only one consecutive potential product state q' such that $q \rightarrow aq'$ where $a \in \Sigma$ is a transition symbol, we can skip

Input : NFA $A_1 = (Q_1, \Sigma, \delta_1, I_1, F_1)$,
NFA $A_2 = (Q_2, \Sigma, \delta_2, I_2, F_2)$
Output: NFA $P = (A_1 \cap A_2) = (Q, \Sigma, \delta, I, F)$ with $L(A_1 \cap A_2) = L(A_1) \cap L(A_2)$

```

1  $Q, \delta, F \leftarrow \emptyset$ 
2  $I \leftarrow I_1 \times I_2$ 
3  $W \leftarrow I$ 
4  $sat \leftarrow False$ 
5  $solved \leftarrow \emptyset$ 
6 while  $W \neq \emptyset$  do
7   picklast  $[q_1, q_2]$  from  $W$ 
8   add  $[q_1, q_2]$  to  $solved$ 
9    $sat \leftarrow \text{areParikhImageFormulaeSatisfiable}([q_1, q_2])$ 
10  if  $sat$  then
11    add  $[q_1, q_2]$  to  $Q$ 
12    if  $q_1 \in F_1$  and  $q_2 \in F_2$  then
13      add  $[q_1, q_2]$  to  $F$ 
14    forall  $a \in \Sigma$  do
15      forall  $q'_1 \in \delta_1(q_1, a), q'_2 \in \delta_2(q_2, a)$  do
16        if  $[q'_1, q'_2] \notin solved$  and  $[q'_1, q'_2] \notin W$  then
17          add  $[q'_1, q'_2]$  to  $W$ 
18        add  $[q'_1, q'_2]$  to  $\delta([q_1, q_2], a)$ 

```

Algorithm 7: Product construction with Parikh image abstraction

computing Parikh image for the state q' as we know for sure Parikh image for this particular product state q' needs to be satisfiable in order to get a satisfiable result for Parikh image for state q . We can add this functionality to our previous algorithm by replacing line 9 with the content of Algorithm 8.

```

1 if not isSkippable( $[q_1, q_2]$ ) then
2   |  $sat \leftarrow \text{areParikhImageFormulaeSatisfiable}([q_1, q_2])$ 
3 else
4   |  $sat \leftarrow \text{True}$ 

```

Algorithm 8: Parikh image computation with skippable states optimization.

3.3.3 Optimization with Incremental SMT Solving

Considering we have to recompute satisfiability Parikh image formula for every potential product state whose Parikh image satisfiability we check, we would appreciate a solution which would allow us to recompute only the clauses which change between two formulae (for two distinct product states) and keep the clauses which remain unchanged from the previous computation to be used in the next computation without the need to recompute them again. Our reduced Parikh image algorithm is designed for such optimization.

Notice that some clauses of Parikh image remain unchanged for the whole automaton, i.e., for every state we compute Parikh image for. Therefore, we can use incremental solving features of SMT, which precompute these clauses only once when Parikh image is first computed⁸. For computing Parikh image for every other state, we make use of these already computed constraints to quicken Parikh image computation.

Assume two original finite automata A and B and state p as a potential product state generated by A and B . The changes of clauses are caused by moving (setting) the corresponding states in both original automata (from which we compute Parikh image formula for satisfiability check)— q in A and s in B —to the current potential product state p as new initial states q_0 and s_0 , respectively, as we proceed further into the automata. We start with the abstract initial states (one for each original automata, q'_0 for A and s'_0 for B).

First, we compute the satisfiability Parikh image formula for $p_0 = (q'_0, s'_0)$, the initial state for A and B are set as q'_0 and s'_0 . If the formula is satisfiable, we generate new potential product states (for example, $p_1 = (q_1, s_1)$ and $p_2 = (q_1, s_2)$). Now we need to check whether to include p_1 and p_2 to the generated product, i.e., check that the Parikh image formula for p_1 and p_2 are satisfiable. Taking p_1 , we set new initial states for both A and B like this: q_1 for A and s_1 for B . Similarly for p_2 , we would set q_1 as a new initial state of A and s_2 as a new initial state of B .

We now need to change every mention of initial states in Parikh image formula because the initial state are different from those we used at the start (q'_0 and s'_0) and for which we already computed the Parikh image formula (state p_1).

Persistent and State Specific Clauses

To present optimization with incremental SMT solving, we split current reduced Parikh image clauses into two groups:

⁸Consequently, computing Parikh image for the first time (for the first state of the given finite automaton) will take longer than for the following product states.

- persistent clauses and
- state specific clauses.

Persistent clauses represent Parikh image clauses which can be precomputed once and used throughout the whole process of working with the given finite automaton. Persistent clauses consist of unchanged clauses of original Parikh image described in 3.3.1:

- clause 2,
- clause 3 and
- clause 4.

State specific clauses are clauses which change with every potential product state tested for satisfiability, and as such have to be constructed and recomputed for every satisfiability check. The whole process of recomputing state specific clauses is the most resource heavy part of the Parikh image computation algorithm, and therefore our goal is to lower the number of state specific clauses as much as possible. The state specific clauses consist of:

- clause 1 as it directly changes the clause according to initial states and
- optionally, if we want to include z_q constraints, clause 3. We would need to recompute z_q constraints for each potential product state too because the clause manipulates with initial states.

As we can see, the majority of clauses can be precomputed for the whole product generation process and only taken into consideration with new state specific clauses. SMT solvers are well optimized to improve their performance by allowing incremental SMT solving.

It is worth to note that the clause 3 manipulates with initial states but the structure of the clause could be reversed to compute connectedness of the automaton in *reversed* order, from the accept states to the initial states, in which case the clause could be reconstructed as a persistent clause dependent on accept states which remain the same (the abstract accept state) for the entire time. This additional optimization is worth inspecting further.

Moreover, SMT solvers can utilize their cache abilities to compute similar, consecutive formulae faster. We can observe how Parikh image satisfiability of successive product states are computed quickly due to minimal changes in formulae which SMT solvers can quickly resolve while using the most of the previously computed formulae constraints.

Algorithm for Incremental SMT solving Using Parikh Image

To implement incremental SMT solving to our current Parikh image algorithm, we need to make the following adjustments.

We need to precompute persistent clauses once for all automata $A_i \in A$. We insert a new line from Algorithm ?? to our algorithm between line 5 and 6. The line contains a call to function `addStatePersistentClauses()` which precomputes all state persistent clauses for all automata $A_i \in A$. Note that the function is called only once, before we enter the *while* loop for iterating over potential product states.

We compute the remaining clauses as normal when we ask for satisfiability of Parikh image formulae for potential product state when we are calling function `areParikhImageFormulaeSatisfiable` on line 9. However, we push the previously precomputed state persistent clauses to the SMT solver stack to preserve them when the current state specific clauses are dropped after the

satisfiability of Parikh image formulae is resolved. For pseudocode of the called function, see Algorithm 9. The function `resolveParikhImageSatisfiability` computes Parikh image formulae and determines their satisfiability, as explained in Section 3.3.1.

```

1 Function areParikhImageFormulaeSatisfiable( $[q_1, \dots, q_i]$ ):
   Data: Potential product state we are determining satisfiability of Parikh image formulae for.
   Result: bool: True if satisfiable, False otherwise.
2   smtSolverPush()
3    $sat \leftarrow \text{resolveParikhImageSatisfiability}([q_1, \dots, q_i])$ 
4   smtSolverPop()
5   return  $sat$ 

```

Algorithm 9: Add state specific clauses to SMT solver for incremental SMT solving optimization.

3.4 Optimization with SMT Solver Timeout

In the case of Parikh images computed with SMT solver, it is easier to find an unsatisfiable counterexample than to prove the formula is satisfiable. Based on our experiments in 4, we use timeout functionalities of SMT solver to quicken the process of resolving satisfiability of potential product states.

We define a maximal amount of time SMT solver can compute satisfiability Parikh image formula for a single product state to resolve its satisfiability for. If SMT solver resolves the satisfiability of such formula before the time runs out, we proceed as normal. However, if the time runs out, the result of satisfiability test is unknown and we must consider such formula might be satisfiable. Then we can, for the sake of efficiency, handle the current formula as a satisfiable one.

This approach resolves the satisfiability of over abstraction of Parikh image satisfiability formulae described above. We prune potential product states whose Parikh image satisfiability can be resolved quickly (within the defined timeout) while allowing the inclusion of some potential product states which are in fact unnecessary to the generated product. Nevertheless, we find pruning capabilities of this optimization satisfactory and the computation time decreases noticeably.

There is a problem with choosing the right timeout time for SMT solver. We say the ideal timeout depends on a structure of finite automata we are working with, their size and complexity. And on how much time we are willing to give to the SMT solver. The timeout time is directly proportional to the results precision and reversely proportional to the scale of over abstraction computed. The cost is that the computation time requirements are directly proportional to SMT timeout time, too.

3.5 Combined Approach of Parikh Image Optimization Enhanced by Length Abstraction

One of the strengths of our optimization algorithms is their high customizability. The different approaches can be combined, easily parallelized and applied on various operations on finite automata. In this section, we present an approach which takes advantage of specific

strengths of our proposed optimization methods while trying to mitigate their weaknesses and utilizes them in a single algorithm.

We introduce a combined algorithm with both length abstraction and Parikh image computation to optimize product construction algorithm.

Input: NFA $A_1 = (Q_1, \Sigma, \delta_1, I_1, F_1)$, NFA $A_2 = (Q_2, \Sigma, \delta_2, I_2, F_2)$

Output: bool: True if satisfiable, False otherwise.

```

1 if not areLengthAbstractionFormulaeSatisfiable( $[q_1, q_2]$ ) then
2   return False
3 addStateSpecificClauses( $[q_1, q_2]$ )
4  $sat \leftarrow$  areParikhImageFormulaeSatisfiable( $[q_1, q_2]$ )
5 if  $sat = \text{true}$  or  $sat = \text{unknown}$  then
6   return True
7 else
8   return False

```

Algorithm 10: Check satisfiability using length abstraction and Parikh image algorithm

3.6 Abstraction of State Language with Mintermization

In this section, we introduce a method of optimizing operations on finite automata using minterms of the given finite automata. Minterm computation abstracts state language of automata in a different approach than which we explored so far, allowing us to follow a diverse set of characteristics about the state language. We can afterwards make use of computed minterms for the automata with other optimization methods introduced in this paper, as well as another optimization approaches.

Foremost, we give an algorithm for minterm computation to compute minterms for the non-empty multiset of input finite automata $A = M_1, M_2, \dots, M_n$, where n equals the number of finite automata, which we desire to execute the required automata operation on. Gained minterms abstract automata state language in such a way we do not lose any information about the former automata, but might create a more concise finite automata which will be easier to work with our other optimization methods and may significantly decrease the computation time required for optimizations such as Parikh image computation.

The general idea is to get sets of transition symbols between two states for all our considered finite automata. Compute minterms from these sets and substitute transition symbols between two states in our automata with corresponding minterms created from these transition symbols. Before all else, let us explain what minterms are and how you can generate them.

Definition 3.6.1 (Minterms) *Given an NFA $M = (Q, \Sigma, \delta, I, F)$, let $\Phi = \{\varphi_1\varphi_2, \dots, \varphi_n\}$ be a finite set of non-empty finite sets of transition symbols $\varphi_i = \{a \mid a \in \Sigma \wedge q \xrightarrow{a} q'\}$, $1 \leq i \leq n$, $q, q' \in Q$ where n equals the number of state pairs (q, q') such that $q \xrightarrow{a} q'$ where $q' \in \delta(q, a)$.*

We call φ_i a *transition set* for the given pair of automaton states q, q' . We denote Ψ or $\text{Minterms}(\Phi)$ as a set of all minterms ψ for the given NFA M such that

$$\text{Minterms}(\Phi) = \left\{ \psi = \bigcap_{1 \leq i \leq n} \psi_i \mid \forall i \in \{1, \dots, n\}. (\psi_i \in \{\varphi, Q \setminus \varphi\}) \wedge \psi \neq \emptyset \right\}.$$

Minterms are computed once, at the beginning of the optimization process for all considered finite automata. We generate so called *minterm tree* with nodes as intersection between sets of transition symbols in case the intersection is non-empty. Each node can have up to two children, representing intersection with the next transition set and its complement, respectively.

When such minterms for the given automaton are computed, we can abstract the state language of the automaton by replacing transitions from the state by their corresponding minterms. We say minterm ψ is created from the set of transition symbols $\varphi \in \Phi$ if φ is used in the intersection defining ψ in its direct form, not as a complement $Q \setminus \varphi$.

Notice that we can compute minterms over multiple NFA, which allows us to use minterms state language abstraction for optimization of operation on those automata.

Let us consider finite automata $M_1 = (\{s_0, s_1, s_2, s_3\}, \Sigma, \delta_1, \{s_0\}, \{s_3\})$ and $M_2 = (\{q_0, q_1, q_2\}, \Sigma, \delta_2, \{q_1\}, \{q_0\})$ over alphabet $\Sigma = \{a, b, c, d\}$ with δ_1 and δ_2 according to Figure 3.9 and Figure 3.10, respectively.




Figure 3.9: Finite automaton M_1 with transitions δ_1 .




Figure 3.10: Finite automaton M_2 with transitions δ_2 .

Figure 3.11: Finite automata M_1 and M_2 used as example automata for mintermization.

The Figure 3.14 depicts how we could mark each transition set in our automata to be used in mintermization process. For example, a transition set φ_1 could be a set of transition symbols from state s_0 to s_1 : $\varphi_i = a, b, d$. Similarly, we mark the remaining transition sets. Now, we can proceed to execute mintermization operations.

If we were to compute minterms for these automata, we would proceed as follows: Starting with the whole alphabet of both automata⁹ at the top of the minterm tree to be generated. Afterwards, we iterate over transition sets. For each transition set φ_i , we compute the intersection of the current minterm tree leaves with:

- the current transition set φ_i and store the result as a left node of this particular tree node,
- the complement of the current transition set $Q \setminus \varphi_i$ and store the result as a right tree node of this particular tree node.

⁹If the automata had non-equal alphabets, we would start with their intersection: $\Sigma = \Sigma_1 \cap \Sigma_2$




Figure 3.12: Finite automaton M_1 with transition sets φ_i .




Figure 3.13: Finite automaton M_2 with transition sets φ_i .

Figure 3.14: Finite automata M_1 and M_2 with marked transition sets used in mintermization.

If the intersection is empty, we omit creating the corresponding child node entirely. In the end, we are left with a complete minterm tree for the given set of transition sets Φ representing the specified finite automata.

The Figure 3.15 illustrates our mintermization process in a diagram.

The acquired minterms are:

$$\Psi = \text{Minterms}(\Phi) = \{\{d\}, \{a, b\}, \{c\}\} = \{\psi_1, \psi_2, \psi_3\}.$$

We can now substitute the former transition sets φ_i for finite automata with the appropriate minterms $\psi_j, 1 \leq j \leq |\Psi|$ which were created from the specific transition sets $\varphi_i \in \Phi$ such that φ_i is used in its direct form (not as a complement) in the process of computing ψ_j . The resulting automata can be seen in Figure 3.18.

Consequently, assuming the previously said, considering we have minterms over alphabet of A , we know that the intersection of two minterms has to be an empty set and that $\forall \psi \in \Psi : \psi \subseteq \varphi, \varphi \in \Phi$ if ψ is created from φ . We make use of this knowledge further.

In the following section, we propose a method of using minterm computation with Parikh image computation optimization. We choose this approach in order to mitigate the disadvantages of Parikh image computation for finite automata, especially those with multitude of transitions between two states varying only in transition symbols, which require considerate time to compute and evaluate. This method proceeds to represent such sets of transitions between two states with a single minterm representing these transitions. We can therefore apply any previously mentioned optimization methods (or any other known optimization method) on such modified automata with minterms as their transition symbols to construct their product without the need to compute, for example, Parikh image for every single transition symbol between two states. We can now compute possibly fewer transitions with the resulting minterms instead.

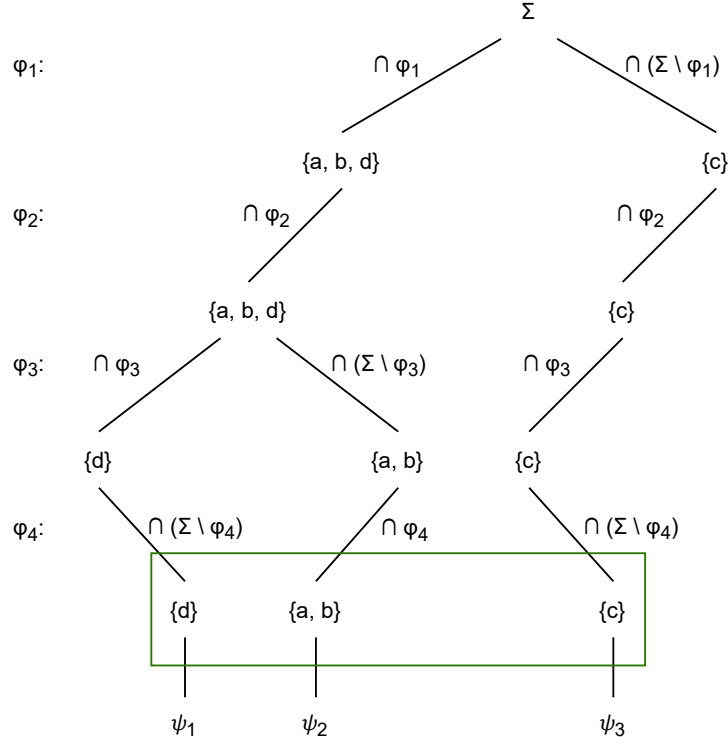


Figure 3.15: Mintermization process executed on example finite automata M_1 and M_2 . We start with the whole alphabet and make our way down through all mintermization sets φ_i , where $1 \leq i \leq n$. For each mintermization set, we compute the intersection of the preceding set with the current mintermization set φ_i . The results are shown in the diagram as the nodes of the tree. When operations on all mintermization sets were executed, the leaves of the tree (indicated by the green square) represent the final minterms for the given mintermization sets Φ over the given alphabet Σ . We denote each minterm ψ_i , where $1 \leq i \leq |\Psi|$ where $|\Psi|$ represents the total number of generated minterms.




Figure 3.16: Finite automaton M_1 with transitions substituted by corresponding minterms $\psi_i \in \Psi$ created from these transition sets. Figure 3.17: Finite automaton M_2 with transitions substituted by corresponding minterms $\psi_i \in \Psi$ created from these transition sets.

Figure 3.18: Finite automata M_1 and M_2 with substituted transitions with minterms in the process of mintermization.

Chapter 4

Experiments and Results

The reference implementation¹ of optimization, written in Python 3, as well as a complete table of all of our experiments and their results and graphs is publicly accessible on a [GitHub repository](#)². There is further explanation of the following graphs as well as additional graphs with description and in-depth analysis of performed experiments.

Test benchmarks used in our experiments were obtained from regular model checking. We have tested various different finite automata and their combinations. We have often used the same automata with their slightly changed variations to simulate real world examples of usually used automata to see how the optimized algorithm reduces the generated state space for certain types of automata with their typical qualities.

We have tested two main aspects:

- First, we have tested the generated state space for emptiness test. That is, whenever we find a solution—accepting state in the intersection, the test ends, and we count the number of generated product states to this moment. If no intersection is found, we end the test when it is certain there is no accepting state and the intersection is indeed empty.
- Second, for the same pair of automata, we have tested the full product construction. Adding new accepting states along the way and comparing generated state spaces in the end for the full product accepting the whole intersection of original automata.

The following graphs show the results for both the emptiness test and full product construction. The graph in Figure 4.1 shows the comparison of product state spaces sizes in basic product construction algorithm and our optimized algorithm considering length abstraction for emptiness test. Sorted in order of increasing product state space size generated by the basic product construction algorithm. The graph in Figure 4.2 shows the same data, only for the full product construction experiment.

Where the length abstraction cannot optimize the product construction, both products have about the same state space size. These results are caused mostly by constructing products of two almost identical automata with only a few states/transitions missing/added which do not affect the accepting runs for recognized languages. There are therefore no branches which can be trimmed—most of the processed states are evaluated as *satisfiable*

¹In the [reference implementation](#), we use [Z3](#) as an SMT solver and automata operations are handled by [our uses modified library Symboliclib](#).

²https://github.com/Adda0/optimizing_automata_product_construction_and_emptiness_test




Figure 4.1: Emptiness test




Figure 4.2: Full product construction

Figure 4.3: Comparison of state space sizes generated by basic and optimized product construction algorithms. Both axes are in logarithmic scale, x-axis showing state space sizes of basic product, y-axis state space sizes of optimized product.

in length abstraction satisfiability check. In full product construction results, if there are nearly no product states to trim, the generated product state space size *explodes* similarly to the basic product construction algorithm—typical for automata with large numbers of transitions from every state causing large numbers of possible accepted lengths, where our algorithm can trim only a few states.

For another automata, the product generated by our algorithm is much smaller. We can see from the graphs that the larger the basic product state space size gets, the higher impact our optimization has on the product state space size. The same holds for the full product construction results. For cases where the intersection is truly empty and accepted lengths differ in both automata, our algorithm stops the process of product construction on the very first tested product state. The basic algorithm continues to create a full product.

We get the best results for automata with practically the same transitions which differ only slightly in final states or a few transitions which affects the accepting runs in the original automata. These changes cause the basic algorithm to generate the product states without realizing most (if not all) product states do not lead to an accepting state. These slight differences in automata (especially in final states) usually also change the length of accepted words. Therefore, our optimization is able to notice these differences and trim most of the product state space, if not the whole product, when no final state can be accessed and the intersection is empty.

In both graphs, we can see the aforementioned quadratic state space explosion for product is nearly not affecting our algorithm in comparison to the basic product construction algorithm. Optimized products are easier to work with and operations on such products require less computational time and memory consumption.

It is worth mentioning that we have neglected the number of generated state space for our lasso automata this whole time. We can use deterministic minimization on original automata to further optimize the generated state space for lasso automata and products. However, we do not need these lasso automata after product construction is complete. Therefore, the lasso automata do not affect how efficiently we work with the generated product. Nevertheless, the number of generated lasso states in the process of deciding the intersection emptiness test matters. For different automata, the generated state space

varies. For state space sizes of lasso automata in our experiments, see [the GitHub repository](#)³.




Figure 4.4: Emptiness test




Figure 4.5: Full product construction

Figure 4.6: Comparison of state space sizes generated by basic and optimized product construction algorithms with sum of states generated for both the final optimized product and lasso automata states generated in the process of the product construction. Both axes are in logarithmic scale, x-axis showing state space sizes of basic product, y-axis state space sizes of optimized product (and optimized product with lasso automata states). The blue dots represent only the optimized product state space sizes (as in Figure 4.3), the orange dots the sum of optimized product state space sizes and the generated lasso states.

As we can see in Figure 4.6, even when counting with lasso automata product state space sizes, the total number of generated states in the whole process of the product construction is usually lower than the basic product state space size. The larger the automata are, the better results we get. It is understandable, that for smaller original automata, whose intersection is computed, the expense of generating lasso automata is significant in comparison with the generated product state space sizes. The larger the original automata get, the lesser the expense of the number of lasso automata states is in comparison with the basic product state space.

Out of all experiments, one weakness of our algorithm is clear—the more final states the original automata have, the more difficult it is to optimize the full product construction using length abstraction. This is caused by the fact that every final state increases the number of accepted different lengths per automaton. Therefore, with automata where out of hundreds or thousands of states nearly every state is a final state too, our optimization algorithm has to consider multiple possible lengths and cannot easily determine which branches will not be accepted by the product automaton.

³https://github.com/Adda0/optimizing_automata_product_construction_and_emptiness_test/tree/master/results




Figure 4.7: Emptiness test




Figure 4.8: Full product construction

Figure 4.9: Stacked comparison of state space sizes generated by basic and optimized product construction algorithms with sum of states generated for both the final optimized product and lasso automata states generated in the process of the product construction. Both axes are in logarithmic scale, x-axis showing state space sizes of basic product (ordered in ascending order), y-axis state space sizes of depicted experiments—because of huge differences in sizes of basic product and optimized product with lasso automata, the largest shown values are set to 6000 and 12000, respectively. Each two columns show a single experiment with our optimized solution as the left (green, red and orange) column—as a sum of all generated states (of optimized product (green) and both lasso automata (red and orange)), and the right blue column as the basic product state space size).

Chapter 5

Conclusion

The most demanding parts of the intersection computation is the generation of product states and transitions of the product automaton. We tried to reduce the size of the generated state space by omitting the states which cannot lead to any accepting state—that is, omitting the *branches* which do not lead to any accepting state—by performing the emptiness test of such states using various state languages abstractions over the original automata such as length abstraction using lasso automata or Parikh image computation based on Parikh’s theorem. Each approach has been experimentally tested and further optimizations to the proposed algorithms were suggested and tested.

According to our experiments, product state space is minimized especially for intersections with huge non-terminating branches or for intersections of automata accepting different lengths of words recognized by the automata languages. Further, for automata with long lines and similar automata varying only slightly from each other. Experiments show our algorithm generates smaller product state spaces for both emptiness test and full product construction, which are two most usually used operations on automata intersection. All our abstractions consider over-approximation of possible products. Therefore, our optimizations are safe to use for any uses resolving operations on finite automata.

We have not encountered similar approaches to product construction optimization using length abstraction or Parikh image computation to compare our results with. It might be worth investing into combining our orthogonal approach with other existing algorithms to see how the generated product state space is affected. We are talking about abstraction techniques such as CEGAR [4] and predicate abstraction [5, 12], IMPACT [17], possibly IC3/PDR [13, 3]. All the above techniques have proven efficient in hardware or software verification, and they can be applied in automata too. First attempts to use these techniques in finite automata problem-solving are based on IC3 [14, 20, 6] and on the interpolation-based approach of McMillan [2, 11].

Bibliography

- [1] ABDULLA, P. A., ATIG, M. F., CHEN, Y., HOLÍK, L., REZINE, A. et al. String Constraints for Verification. In: *CAV*. Springer, 2014, vol. 8559, p. 150–166. Lecture Notes in Computer Science.
- [2] AMLA, N. and MCMILLAN, K. L. Combining Abstraction Refinement and SAT-Based Model Checking. In: *TACAS*. Springer, 2007, vol. 4424, p. 405–419. Lecture Notes in Computer Science.
- [3] BRADLEY, A. R. and MANNA, Z. Checking Safety by Inductive Generalization of Counterexamples to Induction. In: *FMCAD*. IEEE Computer Society, 2007, p. 173–180.
- [4] CLARKE, E. M., GRUMBERG, O., JHA, S., LU, Y. and VEITH, H. Counterexample-Guided Abstraction Refinement. In: *CAV*. Springer, 2000, vol. 1855, p. 154–169. Lecture Notes in Computer Science.
- [5] COLÓN, M. and URIBE, T. E. Generating Finite-State Abstractions of Reactive Systems Using Decision Procedures. In: *CAV*. Springer, 1998, vol. 1427, p. 293–304. Lecture Notes in Computer Science.
- [6] COX, A. and LEASURE, J. Model Checking Regular Language Constraints. *CoRR*. 2017, abs/1708.09073.
- [7] ESPARZA, J. *Automata Theory: An Algorithmic Approach* [online]. 2017 [cit. 2020-10-31]. <https://www7.in.tum.de/~esparza/automatanotes.html>. Available at: <https://www7.in.tum.de/~esparza/automatanotes.html>.
- [8] ESPARZA, J., GANTY, P., KIEFER, S. and LUTTENBERGER, M. *Parikh’s theorem: A simple and direct automaton construction*. June 2011. DOI: 10.1016/j.ipl.2011.03.019.
- [9] FIEDOR, T., HOLÍK, L., JANKU, P., LENGÁL, O. and VOJNAR, T. Lazy Automata Techniques for WS1S. In: *TACAS (1)*. 2017, vol. 10205, p. 407–425. Lecture Notes in Computer Science.
- [10] FIEDOR, T., HOLÍK, L., LENGÁL, O. and VOJNAR, T. Nested antichains for WS1S. *Acta Informatica*. 2019, vol. 56, no. 3, p. 205–228.
- [11] GANGE, G., NAVAS, J. A., STUCKEY, P. J., SØNDERGAARD, H. and SCHACHTE, P. Unbounded Model-Checking with Interpolation for Regular Language Constraints. In: *TACAS*. Springer, 2013, vol. 7795, p. 277–291. Lecture Notes in Computer Science.

- [12] GRAF, S. and SAÏDI, H. Construction of Abstract State Graphs with PVS. In: *CAV*. Springer, 1997, vol. 1254, p. 72–83. Lecture Notes in Computer Science.
- [13] HODER, K. and BJØRNER, N. Generalized Property Directed Reachability. In: *SAT*. Springer, 2012, vol. 7317, p. 157–171. Lecture Notes in Computer Science.
- [14] HOLÍK, L., JANKU, P., LIN, A. W., RÜMMER, P. and VOJNAR, T. String constraints with concatenation and transducers solved efficiently. *Proc. ACM Program. Lang.* 2018, vol. 2, POPL, p. 4:1–4:32.
- [15] KOZEN, D. C. Parikh’s Theorem. In: *Automata and Computability*. Berlin, Heidelberg: Springer Berlin Heidelberg, 1977, p. 201–205. ISBN 978-3-642-85706-5. Available at: https://doi.org/10.1007/978-3-642-85706-5_35.
- [16] LIN, A. W. and BARCELÓ, P. String solving with word equations and transducers: towards a logic for analysing mutation XSS. In: *POPL*. ACM, 2016, p. 123–136.
- [17] MCMILLAN, K. L. Lazy Abstraction with Interpolants. In: *CAV*. Springer, 2006, vol. 4144, p. 123–136. Lecture Notes in Computer Science.
- [18] SIEGEL, S. F. and YAN, Y. Action-Based Model Checking: Logic, Automata, and Reduction. In: *CAV (2)*. Springer, 2020, vol. 12225, p. 77–100. Lecture Notes in Computer Science.
- [19] SIPSER, M. *Introduction to the Theory of Computation*. 3rd ed. Cengage Learning, 2013. ISBN 13: 978-1-133-18779-0.
- [20] WANG, H., TSAI, T., LIN, C., YU, F. and JIANG, J. R. String Analysis via Automata Manipulation with Logic Circuit Representation. In: *CAV (1)*. Springer, 2016, vol. 9779, p. 241–260. Lecture Notes in Computer Science.