BRNO UNIVERSITY OF TECHNOLOGY
VYSOKÉ UČENÍ TECHNICKÉ V BRNĚ

FACULTY OF INFORMATION TECHNOLOGY
FAKULTA INFORMAČNÍCH TECHNOLOGIÍ

DEPARTMENT OF INTELLIGENT SYSTEMS
ÚSTAV INTELIGENTNÍCH SYSTÉMŮ

ABSTRACTION OF STATE LANGUAGES IN AUTOMATA ALGORITHMS

ABSTRAKCE JAZYKŮ STAVŮ V AUTOMATOVÝCH ALGORITMECH

BACHELOR'S THESIS
BAKALÁŘSKÁ PRÁCE

AUTHOR   DAVID CHOCHOLATÝ
AUTOR PRÁCE

SUPERVISOR   doc. Mgr. LUKÁŠ HOLÍK, Ph.D.
VEDOUCÍ PRÁCE

BRNO 2022

Bachelor's Thesis Specification

25013

Student: **Chocholatý David**

Programme: Information Technology

Title: **Abstraction of State Languages in Automata Algorithms**

Category: Algorithms and Data Structures

Assignment:

The goal is to explore possibilities of using various abstractions of automata languages in optimisation of automata algorithms.

We will start with abstracting languages of states to sets of possible word lengths and to Parikh images, represented as semi-linear sets, and exploring options of using them to optimize the construction of synchronous product of automata by pruning pairs of states with incompatible abstractions. We will then continue either towards optimisation of these techniques or towards searching for alternatives or more advanced versions.

1. Study the literature on automata and Parikh images, familiarise yourself with technology of SMT solvers.
2. Implement automata product construction optimised with length abstraction and Parikh image abstraction of state languages.
3. Study the state pruning capabilities of these abstractions.
4. If the pruning capabilities prove promising, study possibilities of their efficient implementation and compare it with standard synchronous product construction on provided data.
5. Otherwise continue the research by elaborating on the principle of using state language abstractions to optimize automata constructions.

Recommended literature:

- Javier Esparza, Pierre Ganty, Stefan Kiefer, Michael Luttenberger: Parikh's Theorem: A simple and direct construction. CoRR abs/1006.3825, 2010.

Requirements for the first semester:

- Items 1 and 2.

Detailed formal requirements can be found at https://www.fit.vut.cz/study/theses/

Supervisor: **Holík Lukáš, doc. Mgr., Ph.D.**

Head of Department: Hanáček Petr, doc. Dr. Ing.

Beginning of work: November 1, 2021

Submission deadline: May 11, 2022

Approval date: November 3, 2021

Abstract

We explore possibilities of using various abstractions of finite automata languages in optimization of automata algorithms used in automata reasoning. We focus on abstracting languages of states to sets of accepted lengths of word or Parikh images, represented as semi-linear sets, and explore options of using them to optimize automata constructions by pruning states based on abstractions of their languages. We propose several abstractions and work on optimizing their performance. We use two common finite automata problems, synchronous product construction and deciding the emptiness of finite automata intersection, as benchmark problems on which we test our optimizations. Nevertheless, our abstractions are applicable on many other typical automata operations, e.g., complement generation etc. Our experiments show that the proposed optimizations reduce generated state space for both benchmark problems substantially.

Abstrakt

Prověřujeme možnosti použití různých abstrakcí jazyků konečných automatů pro optimalizaci automatových algoritmů používaných pro rozhodování založeném na automatech. Zajímáme se o abstrakci jazyků stavů na množiny přijímaných délek slov nebo Parikovy obrazy, reprezentované jako semi-lineární množiny, a zkoumáme možnosti jejich využití k optimalizaci automatových konstrukcí odstraňováním stavů založené na abstrakcích jejich jazyků. Předvádíme několik abstrakcí a pracujeme na optimalizaci jejich výkonu. Používáme dva běžné automatové problémy, synchronní produkt konstrukci a rozhodování prádnosti průniku konečných automatů, jako operace pro experimentální vyhodnocení, na kterých testujeme naše optimalizace. Naše abstrakce jsou nicméně aplikovatelné na mnohé další typické automatové operace, například generaci doplňku aj. Provedené experimenty ukazují, že navrhované optimalizace podstatně zmenšují generovaný stavový prostor pro oba testované problémy.

Keywords

Klíčová slova

Reference

CHOCHOLATÝ, David. *Abstraction of State Languages in Automata Algorithms*. Brno, 2022. Bachelor's thesis. Brno University of Technology, Faculty of Information Technology. Supervisor doc. Mgr. Lukáš Holík, Ph.D.

Rozšířený abstrakt

Konečné automaty nachází mnohá využití v různých oblastech výpočetní teorie, zejména v oblasti rozhodování založeném na automatech (model checking, string solving a analýza, WS1S). Přestože jsou konečné automaty konceptuálně jednoduché, často s nimi potřebujeme provádět operace, které jsou výpočetně drahé a generují rozsáhlý stavový prostor, jehož mnohé části jsou nadbytečné.

V této práci zkoumáme možnosti použití různých abstrakcí jazyků stavů automatů pro optimalizaci takových automatových algoritmů. Pomocí vhodných abstrakcí se snažíme předpovědět, které stavy výsledného automatu jsou nepotřebné, a mohou proto být odstraněny z generovaného stavového prostoru bez narušení jazyku výsledného automatu, pokud jsou získané abstrakce navzájem nekompatibilní.

Pro demonstraci našich abstrakcí jsme se rozhodli použít operaci průniku konečných automatů prováděnou synchronní produkt konstrukcí a test prázdnosti průniku automatů. Naše předvedené abstrakce jsou však navrženy tak, aby byly aplikovatelné na širokou škálu automatových operací (například konstrukci doplňku aj.). Význam naší práce proto přesahuje samotnou optimalizaci produkt konstrukce automatů. Všechny navrhované abstrakce s jejich inverzními funkcemi navíc tvoří *Galois connection*, tedy popisují nadabstrakci jazyků stavů. Díky tomu není nebezpečí, že bychom při odstraňování stavů s nekompatibilními abstrakcemi nechtěně odstranili i stavy důležité pro popis jazyka přijímaného generovaným automatem.

Při konstrukci průniku automatů dochází k tzv. stavové explozi, kdy jsou generovány rozsáhlé části stavového protoru, které tvoří neukončující stavy, ze kterých nebude dosažitelný žádný koncový stav ve výsledném produktu. Naše optimalizace sestává z kontroly kompatibility abstrakcí jazyků stavů pro stavy, ze kterých se skládá daný produkto-stav, za běhu produkt konstrukce. Pokud určíme abstrakce jako nekompatibilní, můžeme bezpečně takový produkto-stav odstranit. Výhodou našich abstrakcí je, že stavový prostor zmenšují již při generaci výsledného automatu. Některé stavy tak nebude třeba vůbec ani generovat, pokud všechny jejich předchůdci budou odstraněni. Naproti tomu u naivní produkt konstrukce musíme nejdříve vygenerovat celý automat, než můžeme rozhodovat o kompatibilitě jazyků vstupních automatů.

Mezi zkoumané abstrakce jazyků stavů patří délková abstrakce a abstrakce Parikovými obrazy. Dále zkoumáme možnosti optimalizace těchto abstrakcí či předzpracování vstupních automatů, například pomocí mintermizace automatů.

Délková abstrakce tvoří nadaproximaci jazyka stavů na lineární množiny možných délek slov přijímaných jazykem pomocí lineárních délkových formulí. Aby daný produkto-stav patřil do průniku, přijímané délky slov stavů ve vstupních automatech si musí odpovídat, tedy formule popisující délkovou abstrakci musí být splnitelné zároveň. V opačném případě jazyky stavů nepřijímají stejný jazyk (délky přijímaných slov se liší) a jejich průnik je prázdný. Takové produkto-stavy mohou být odstraněny z generovaného stavového prostoru a jejich následníci nemusí být generováni.

Délky slov modelujeme pomocí tzv. laso automatů přijímajících nadmnožinu jazyka vstupních automatů přijímající slova o všech délkách slov přijímaných vstupními automaty. Vzájemnou splnitelnost délkových abstrakcí v podobě délkových formulí sestavených z laso automatů ověřujeme zadáním příkazu pro SMT solver, nicméně můžeme optimalizovat otázku splnitelnosti délkových formulí nahrazením SMT solveru za matematický výpočet založený na vlastnostech lineární kongruence, který je schopný rychle a efektivně rozhodnout o splnitelnosti délkových formulí.

Abstrakce Parikovými obrazy definuje semi-lineární množiny založené na Parikově teorému abstrahující jazyky stavů na počty výskytů symbolů na přechodech bez závislosti na jejich umístění v přijímaném slově pomocí semi-lineárních formulí Parikových obrazů. Za nekompatibilní abstrakce považujeme takové, kde si neodpovídají počty použitých symbolů jazyků stavů pro daný produkto-stav. Tedy, pokud jsou formule Parikových obrazů navzájem nesplnitelné, můžeme opět odstranit daný produkto-stav z generovaného stavového prostoru.

Abstrakci Parikovými obrazy je možné nadále optimalizovat další redukcí Parikových obrazů či inkrementálním SMT výpočtem, který umožňuje předpočítat společné části formulí jednou a využívat výsledky předchozího výpočtu po celou konstrukce produktu. Nadále můžeme zavést *timeout* pro předčasné ukončení rozhodování splnitelnosti formulí Parikových obrazů.

Obě abstrakce mohou využít optimalizace přeskočitelných produkto-stavů, kdy není třeba vyhodnocovat splnitelnost formulí abstrakcí, pokud daný produkto-stav byl vytvořen z produkto-stavu generujícího pouze tento jediný následující produkto-stav. Tedy, aby měl předcházející produkto-stav kompatibilní abstrakce jazyků stavů vstupních automatů, musí využívat aktuálního produkto-stavu pro dosažení koncového stavu, a proto musí nutně i abstrakce jazyků stavů pro tento následující produkto-stav být navzájem kompatibilní.

Naše abstrakce jsme navrhli tak, aby tvořily obecný a samostatný popis jazyků stavů, což umožňuje abstrakce volitelně kombinovat, rozšiřovat o další abstrakce či optimalizační techniky, a tím využít výhod každé abstrakce, zatímco minimalizujeme dopad nevýhod daných abstrakcí. Tím umožňujeme využívat naše optimalizace pro širokou oblast problémů řešených konečnými automaty. Přístup za běhu řešených abstrakcí jazyků stavů taktéž umožňuje operace paralelizovat nebo vhodně rozdělit na podproblémy.

Provedli jsme experimentální vyhodnocení navrhovaných abstrakcí optimalizujících konstrukci průniku. Podle provedených experimentů můžeme soudit, že navrhované abstrakce mají předpokládané optimalizační schopnosti a zmenšují generovaný stavový prostor i lépe rozhodují test prázdnosti průniku automatů než naivní přístupy konstrukce produktu.

Délková abstrakce je rychlá a jednoduchá, její optimalizační síla je však nižší než u abstrakce Parikovými obrazy. Délková abstrakce výborně optimalizuje produkty s dlouhými linkami stavů, může mít však potíže s odstraňováním stavů v hustě propletené síti přechodů. Abstrakce Parikovými obrazy je velmi přesná. Skvěle optimalizuje generovaný produkt, ovšem výpočet vzájemné splnitelnosti formulí Parikových obrazů je pro SMT solver náročný a časově drahý. Můžeme si tedy zvolit, jestli chceme dosáhnout rychlého, i když možná méně důkladného zmenšení stavového prostoru; přesné, ale výpočetně náročnější minimalizace průniku; případně vhodné kombinace těchto vlastností.

Abstraction of State Languages
in Automata Algorithms

Declaration

I hereby declare that this Bachelor's thesis was prepared as an original work by the author under the supervision of doc. Mgr. Lukáš Holík, Ph.D. I have listed all the literary sources, publications and other sources, which were used during the preparation of this thesis.

. . . .
David Chocholatý
May 11, 2022

Acknowledgements

Contents

Chapter 1

Introduction

Finite automata are a well-known model of computational theory used in many areas. Finite automata are commonly used in automata reasoning (e.g., in model checking [28], string solving and analysis [26] or WS1S [14, 15]).

Finite automata are conceptually straightforward. However, operations on finite automata are often expensive: have high complexity, require extensive computational time and generate vast state space.

Our goal is to find different heuristics for optimizing several typical problems connected to finite automata. We study possibilities of using various abstractions of languages of finite automata states in optimization of automata algorithms. We abstract languages of states to sets of lengths of words in state languages and to Parikh images, represented as semi-linear sets, and explore options of using them to optimize the automata constructions by pruning states whose language abstractions represent an empty language. We work on optimizing performance of these abstractions. Moreover, besides optimization techniques specific for concrete state language abstractions, we also consider a general technique of mintermization to allow using our abstractions on additional structures such as regular expressions represented as finite automata and to optimize our abstractions further.

The idea of using abstraction in automata problem-solving is not new, but it is not properly explored either. There were first attempts of using abstraction techniques in automata such as alternating automata [17] or abstract regular model checking [5], both using techniques similar to a general predicate abstraction [8, 18] and CEGAR [7].

We want to optimize operations on finite automata which take lots of computational time and generate vast state space. We are considering operations such as product construction, determinization of complement construction, minimization or determinization and inclusion test. Furthermore, we want to create state language abstractions which can work for different automata structures: operations on transducers, operations with alternating automata such as its emptiness or a conversion of an alternating automaton to its NFA representation, conversion of finite automata to flat automata, etc.

We focus on the construction of finite automata intersection generated by the synchronous product construction. We consider two common forms of this finite automata operation as our benchmark problems on which we test our optimizations:

- first, construction of the intersection automaton by the synchronous product construction, which means completion of the entire product construction, and

- testing the emptiness of finite automata intersection (emptiness problem) which asks whether the language of the product is empty. Here, it is not always necessary to

construct the entire product (or parts of the product) to resolve the emptiness of the intersection.

Nevertheless, even if our optimizations are introduced on product construction and emptiness problem, our discoveries have wider impact and are in some form applicable on many typical automata operations.

The intersection of finite automata combines the original states from the individual automata to tuples called product states in the generated state space by finding corresponding transitions with the same symbols. Every product state represents an intersection of languages of the corresponding states in the original automata. The synchronous product construction is computationally costly: for two finite automata, the generated product state space can increase quadratically to the number of input finite automata states (number of states in one finite automaton times number of states in the second finite automaton) and transitions. And, for multiple finite automata, exponentially to the number of used finite automata. However, there are often large parts of the generated state space which cannot accept any words (no final states can be reached from these states), yet are still generated[1]. Therefore, it is important to have a decent algorithm to minimize the generated product state space as much as possible.

In our optimizations, we try to identify which generated product states cannot lead to any accepting state or are successive only to such states. When state language abstractions of states in product state are not compatible—the original languages of the corresponding states cannot accept the same words—we can omit such product state and all their potential successive states, pruning the generated state space.

We start with an optimization using length abstraction of state languages. For each state, we construct a so-called lasso automata. It is used to compute a semi-linear formula which codes the lengths of the words in the languages of current states (we call them accepted lengths). We use SMT solver to resolve satisfiability of these formulae. When accepted lengths of states in the product states are not compatible (their formulae are not satisfiable), their languages have no common words. There is no path from the product state leading to the accepting product state. We can prune such product states. Consequently, this removes the need to even consider their potential successive states.

Even though there still might be states which do not lead to any final state in the final product, this simple optimization often trims substantial parts of the state space. Length abstraction can also be implemented simply and efficiently. However, sometimes the abstraction is too coarse. For instance, it cannot detect unnecessary product states for finite automata with rich alphabets, since their states accept a multitude of word lengths.

For that reason, we investigate a finer state language abstraction which uses Parikh images of state languages. Parikh image of a word tells us how many times each symbol occurs in the word[2]. Parikh image of a language is a semi-linear formula describing the relation between the number of symbol occurrences in words in a language. In contrast to the length abstraction, it contains additional information about the numbers of symbols in words. We can more precisely identify unnecessary state space by determining compatibility of Parikh image abstractions. To solve satisfiability of their formulae, we use SMT solver. However, the Parikh image computation is expensive. There is a trade-off between the precision of the Parikh image abstraction and its cost.

---

[1]The generated product state space sometimes *explodes*.

[2]A function which assigns each transition symbol a number of occurrences in a word.

Generating smaller state space using our Parikh image optimization can improve computation time for the product generation in case substantial parts of the state space are pruned. Moreover, it is even enough to decide the emptiness of the intersection on the initial product state immediately in many cases.

An important part of this work is researching optimizations of the specific state language abstractions to make their usage efficient. For both abstractions, we find a way to skip evaluation of state language abstractions for product states in long *lines* (linear non-branching sequences of states). If a product state with compatible accepted lengths generates a single sequence of states, all states in the line have compatible length abstractions.

For length abstraction, we reduce lasso automata generation for each state to a single expanding lasso automaton for the whole finite automaton. We also efficiently evaluate length abstractions without SMT solver by resolving satisfiability of their formulae with a special construction using linear congruences.

For Parikh image computation, we remove parts of Parikh image formula, which can reduce its precision, but it occurs that it has no impact on pruning capabilities on our benchmark automata. The formulae are large. However, extensive parts of them remain unchanged for different product states. We utilize incremental SMT solving to precompute the common parts and for each state recompute only the remainder of the formulae. This speeds up the evaluation of compatibility of Parikh image abstractions. If resolving satisfiability of formulae takes too long, we can introduce a timeout for SMT solver to stop the computation and not prune the product state.

We also consider combinations of our abstractions, particularly we experiment with computing cheap length abstraction first and computing the Parikh images only when length abstraction fails to prune product states.

Further, we use mintermization for intersection of finite automata as a different approach to processing the initial automata before applying other optimizations. We compute minterms, which can be used instead of transition symbols while retaining all information about the automata to compute Parikh images and other optimization abstractions faster.

We implement the proposed abstractions and evaluate their impact on the emptiness problem and the product construction experimentally. We experiment with a benchmark containing a set of different finite automata obtained from runs of a regular model checking tool on verification of pointer programs and parametric protocols created in [4] based on a method of abstract regular model checking from [5]. We generate products of various combinations of these finite automata and test emptiness of product languages to solve emptiness problem of their intersections. We focus on the number of trimmed product states and their nature, their position in the product or other significant properties. For certain types of automata, our optimizations work really well. Parikh image abstraction usually trims vast state spaces where length abstraction cannot prune everything and unoptimized product state space explodes (e.g., from 20000 to 10 product states). In addition, the abstractions are sometimes successful at immediately stopping product construction on the first initial product state if the intersection is empty, while, in some cases, product construction would take hours (in one case, more than 7 hours, compared to 1 minute with Parikh image abstraction).

The contribution of this work can be summarized as follows:

1. heuristics trimming the generated state space of finite automata operations based on abstractions of specific state languages: length abstraction and Parikh image computation; or general approaches as mintermization,

2. optimizations for explored state language abstractions:

   - skipping evaluation of state language abstractions for some product states for sequences of product states in long lines,

3. optimizations specific for length abstraction:

   - generating a single lasso automaton for the whole finite automaton,
   - efficient evaluation of length abstraction without SMT solver,

4. optimizations specific for Parikh image abstraction:

   - reduced Parikh image to resolve satisfiability of Parikh image formulae faster,
   - resolving Parikh images with incremental SMT solving,
   - resolving Parikh images with a timeout for SMT solver,

5. combination of state language abstractions to optimize automata problems, and

6. implementation and experimental evaluation of said heuristics and their optimizations.

Chapter 2

Preliminaries

Let us clarify a few definitions and terms often used throughout this paper. The following definitions are mostly adapted from [12] or [29].

   Alphabet is a finite, non-empty set denoted by Σ. Elements of an alphabet are called *symbols* or *letters*. A finite, possibly empty, sequence of symbols over an alphabet is a *word* w from the set of all words Σ^* over an alphabet Σ.

Definition 2.0.1 (Deterministic finite automaton)
A deterministic finite automaton (DFA) is a 5-tuple $A = (Q, \Sigma, \delta, I, F)$, where:

- *Q is a non-empty **set of states**,*

- *Σ is an **input alphabet**,*

- *δ is a **transition function**: $Q \times \Sigma \to Q$,*

- *$I \in Q$ is an **initial state**, and*

- *$F \subseteq Q$ is a **set of final (accepting) states**.*

   A *run* of A on input $a_0 a_1 a_2 \dots a_{n-1}$ is a sequence $q_0 \xrightarrow{a_0} q_1 \xrightarrow{a_1} q_2 \xrightarrow{a_3} \dots \xrightarrow{a_{n-1}} q_n$, such that $q_i \in Q$ for $0 \le i \le n$, $q_0 = I$ and $\delta(q_i, a_i) = q_{i+1}$ for $0 \le i \le n-1$. A run is *accepting* if $q_n \in F$. A *accepts* a word $w \in \Sigma^*$ if A has an accepting run on input w. A *language* recognized by A is a set $L(A) = \{w \in \Sigma^* \mid w \text{ is accepted by } A\}$. A single *transition* from δ is denoted as $q \xrightarrow{a} q'$ if $q' \in \delta(q, a)$ and means *one can get from state q to state q' with a transition symbol a*. For every state, DFA has at most one transition for a given symbol. Consequently, DFA has exactly one run on a given word from initial state to one of the accepting states (or non-terminating states[1] in case the word is not accepted by the automaton at all).

Definition 2.0.2 (Non-deterministic finite automaton)
A non-deterministic finite automaton (NFA) is a 5-tuple $A = (Q, \Sigma, \delta, I, F)$, where Q, Σ and F are as for DFA and:

- *δ is a **transition relation**: $\delta : Q \times \Sigma \to P(Q)$, where $P(Q) = \{R \mid R \subseteq Q\}$ is a set of subsets of Q, and*

- *$I = \{q \mid q \in Q\}$ is a non-empty **set of initial states**.*

---

[1]No accepting state is accessible from them.

For every state and its transition symbol, $P(Q) \in \delta(q, a)$ is a singleton. For example, $\delta(q_1, a) = \{q_1, q_2\}$.

Two finite automata A_1 and A_2 are said to be *equivalent* when both accept the same language: $L(A_1) = L(A_2)$.

For every NFA A exists a corresponding equivalent DFA A'. *Determinization* is a process of converting such NFA to DFA.

Definition 2.0.3 (Powerset (subset) construction)

The powerset construction is a method for creating a corresponding deterministic finite automaton from its equivalent non-deterministic finite automaton. Produces finite automaton A', where $Q' = 2^Q$, $F' = \{S \in Q' | S \cap F \neq \varnothing\}$, $I' = I$ and for $S \in Q' : \delta'(S, a) = \bigcup_{s \in S} \delta(s, a)$.

Definition 2.0.4 (Product construction)

Given two NFAs $A_1 = (Q_1, \Sigma, \delta_1, I_1, F_1)$ and $A_2 = (Q_2, \Sigma, \delta_2, I_2, F_2)$ over the same alphabet Σ, operations on A_1 and A_2 yield a result—a product A as a 5-tuple deterministic finite automaton $A = (Q, \Sigma, \delta, I, F)$ where:

- *$Q = Q_1 \times Q_2$,*

- *$\delta : Q \times \Sigma \rightarrow P(Q)$,*

- *$I = I_1 \times I_2$, and*

- *$F = F_1 \times F_2$.*

δ is described as $([q_1, q_2], a) = \delta_1(q_1, a) \times \delta_2(q_2, a)$. For pairs of states q_1 and q_2 from A_1 and A_2, respectively, and a common transition symbol a of transitions $q_1' \in \delta_1(q_1, a)$ and $q_2' \in \delta_2(q_2, a)$, we denote a single product transition as $[q_1, q_2] \xrightarrow{a} [q_1', q_2']$, where $[q_1', q_2'] \in \delta([q_1, q_2], a)$ for the corresponding states $[q_1, q_2]$ and $[q_1', q_2']$ in A are called *product states*.

Focusing on an *intersection* of finite automata, the product construction tells that $L(A) = L(A_1) \cap L(A_2)$. Finally, we test the *emptiness* of the intersection. Given A_1 and A_2, emptiness test asks whether the language of the product is empty: $L(A_1 \cap A_2) = \varnothing$.

We work with an unoptimized product construction in Algorithm 1.

Definition 2.0.5 (Galois Connection)

Galois connection is a quadruple $\pi = (\mathcal{P}, \alpha, \gamma, \mathcal{Q})$ such that:

- *$\mathcal{P} = \langle P, \leq \rangle$ and $\mathcal{Q} = \langle Q, \sqsubseteq \rangle$ are partially ordered sets (posets) and*

- *abstraction function $\alpha : P \rightarrow Q$ and concretization function $\gamma : Q \rightarrow P$ inverse to α. $\forall p \in P$ and $\forall q \in Q$:*

$$p \leq \gamma(q) \Leftrightarrow \alpha(p) \sqsubseteq q.$$

In the terminology of abstract interpretation, P is a *concrete domain* and Q is an *abstract domain*. If α and γ functions form a Galois connection, $\forall p \in P \big(p \leq \gamma(\alpha(p))\big)$. That is, the abstraction may only over-approximate the concrete semantics.

Input : NFA $A_1 = (Q_1, \Sigma, \delta_1, I_1, F_1)$, NFA $A_2 = (Q_2, \Sigma, \delta_2, I_2, F_2)$
Output: NFA $A = (A_1 \cap A_2) = (Q, \Sigma, \delta, I, F)$ with $L(A_1 \cap A_2) = L(A_1) \cap L(A_2)$

1 $Q, \delta, F \leftarrow \varnothing$
2 $I \leftarrow I_1 \times I_2$
3 $W \leftarrow I$
4 **while** $W \neq \varnothing$ **do**
5   **pick** $[q_1, q_2]$ **from** W
6   **add** $[q_1, q_2]$ **to** Q
7   **if** $q_1 \in F_1$ *and* $q_2 \in F_2$ **then**
8   **add** $[q_1, q_2]$ **to** F
9   **forall** $a \in \Sigma$ **do**
10   **forall** $q_1' \in \delta_1(q_1, a), q_2' \in \delta_2(q_2, a)$ **do**
11   **if** $[q_1', q_2'] \notin Q$ **then**
12   **add** $[q_1', q_2']$ **to** W
13   **add** $[q_1', q_2']$ **to** $\delta([q_1, q_2], a)$

Algorithm 1: Classic unoptimized product construction used by our state language abstractions to optimize the generated product state space by deciding the compatibility of state language abstractions.

Chapter 3

State Language Abstractions

In this chapter, we introduce several state language abstractions, presented on product construction and deciding the emptiness problem.

When constructing a product, a considerate number of product states are non-terminating and thus unnecessary. Moreover, the whole product must be constructed before we can determine whether the automata intersection is empty. We want to minimize the number of generated product states when resolving the product construction of automata intersection and deciding the emptiness of the intersection.

We try to guess which product states do not lead to any final states and consequently can be omitted, and the following states do not need to be generated at all. Our optimizations decide the emptiness of parts of the product (or the whole product) already in the process of generating the product (on the fly). We can thus prune non-terminating states before they are added to the product and omit extensive product state space before even considering it in the classic product construction. We achieve this by computing state language abstractions for each state the generated product state consists of and deciding the compatibility of these abstractions.

Our product construction optimizations are applicable on two and more automata, but for the ease of explanation, we consider only two automata. In the following, we will define an abstraction of languages of the states q, $\alpha(q)$. We define two kinds of abstractions: length abstraction $\alpha^{LA}(q)$ and Parikh image abstraction $\alpha^{PI}(q)$. These abstractions represent formulae in first-order predicate logic. Both our $\alpha^{LA}(q)$ and $\alpha^{PI}(q)$ together with their inverse functions form a Galois connection. Hence, they are an over-approximation of state language of q.

For a product state $p = [q_1, q_2]$ of the product P, we use abstractions of languages of states $\alpha(q_1)$ and $\alpha(q_2)$ to quickly detect whether p has an empty language. That can be achieved by checking whether $\alpha(q_1)$ and $\alpha(q_2)$ are compatible. If they are incompatible, product language is empty and p can be pruned (there is no run from p to any final state). Therefore, the optimized product language is the same as the unoptimized product language.

3.1 Length Abstraction of State Languages

In this section, we discuss length abstraction $\alpha^{LA}(q)$. Length abstraction looks at lengths of words accepted by the state language, creating a set of accepted lengths. In the following, we first discuss the basic principle of length abstraction. Later, we propose efficiency

optimizations for length abstraction. To start with, we introduce lasso automata as a finite automata representation of length abstraction.

3.1.1   Length Abstraction Represented by Lasso Automata

Length abstraction over-approximates the language of q by considering only the accepted lengths of words. This is, if a length of a word does not belong to the length abstraction of q, it cannot be accepted by state language of q, either.

Computing length abstraction over the languages of finite automata states is accomplished using lasso automata (LSA, handle and loop automata)—deterministic finite automata with a unary alphabet (similar as in [1]). They consist of a *handle* (a sequence of states from the initial state) and a *loop* (resolving the cycles in the original automaton) resembling a lasso with a few final states representing the accepted word lengths.

You can create a lasso automaton for a state q, $lsa(q)$, by taking the finite automaton A, $q \in Q_A$, setting $I_A = \{q\}$, considering all transition symbols as a single transition symbol and determinizing the result with subset construction. $lsa(q)$ is an automaton accepting every length of any word in state language of q. Consequently, it is easy to compute semi-linear set (formulae in the form of a disjunction of linear equations) for the accepted lengths of words, which can be efficiently evaluated. We are computing length formulae for individual states in the product state, checking their satisfiability and constructing only those product states for which the length abstraction formulae are satisfiable.

The length abstraction formulae are generated from $lsa(q)$. For every state q, we get one or more existentially quantified formulae φ in Presburger arithmetic describing language abstracting $\alpha^{LA}(q)$ in the form

$$\varphi : \exists k(|w| = h + l \cdot k)$$

where $|w|$ is a length of a recognized word, h is the length of a handle to a certain final state f, and l is the length of a loop to return to f going through the loop. k is the number of cycles through the loop states until a word ends in f. When multiple depicted formulae are created (because there are more final states or different accepting runs for a single final state in LSA resulting in multiple accepted lengths), we append these formulae with *logical or*:

$$\alpha^{LA} : \exists k(\varphi_1 \vee \ldots \vee \varphi_n)$$

where n is a number of generated φ.

Running Example   We demonstrate the construction of $\alpha^{LA}(q_0)$ for initial state q_0 of the following NFA $A_1 = (\{q_0, q_1, q_2, q_3, q_4, q_5\}, \{0, 1\}, \delta_1, \{q_0\}, \{q_4\})$ where transition relation δ_1 is depicted in Figure 3.1. That is, we construct $lsa(q_0)$. We will continue using A_1 throughout the section.

A_1 is a non-deterministic finite automaton (see state q_1) and uses multiple input symbols. Due to the fact we work with only recognized word lengths, we can substitute the automaton alphabet with a unary alphabet of a single input symbol $*$[1]. See the obtained finite automaton A_1' in Figure 3.2.

---

[1]Even though we do not actually need any particular input symbol, we use $*$ here as an example to depict the process. In general, all we need to know is that there is a transition between two states. The specific transition symbols are not significant for our length abstraction.

Figure 3.1: Non-deterministic finite automaton A_1.

Figure 3.2: Non-deterministic finite automaton A_1': A_1 with unified transition symbols.

Then, we can generate $lsa(q_0)$, which is its deterministic equivalent. For the final $lsa(q_0)$, generated from A_1' by its product construction determinization, see Figure 3.3. Lasso automaton for q_0 accepts any words of lengths of words recognized by state language of q_0.


Figure 3.3: Lasso automaton $lsa(q_0)$ for the original NFA A_1 generated from A_1' by its determination with subset construction.

3.1.2 Single Lasso Automaton for Each Original Automaton

When we are constructing a product of finite automata A and B, we do not want to regenerate $lsa(q_A)$ and $lsa(q_B)$ for each product state $p = [q_A, q_B]$. This is inefficient. Due to the nature of LSAs, the successive product states $p' = [q_A', q_B']$ generate LSAs very similar to the LSAs for p. We can construct one summary lasso automaton for the whole finite automaton A, $LSA(A)$, which contains all the lasso automata $lsa(q_A)$ for all states in A. Similarly for B.

$LSA(A)$ is a lasso automaton created as a union of all $lsa(q_A)$ for each $q_A \in Q_A$. That is, $LSA(A)$ is a union of all lasso states, transitions, final and initial states for each q_A. As a result, α^{LA} for product construction generates only one LSA for each finite automaton (possibly with multiple loops and/or multiple handles). All singleton states in $LSA(A)$ can be initial states. The initial state changes to which state q_A we want $lsa(q_A)$ for.

To generate only the necessary $lsa(q_A)$, we can construct $LSA(A)$ gradually, state by state, for only the currently required q_A. We generate $lsa(q_A)$ as the first part of $LSA(A)$. If we need $lsa(q'_A)$ later, we extend $lsa(q_A)$ with $lsa(q'_A)$ by union of both LSAs. When the new LSA state l_A is not already present in $LSA(A)$, we add l_A to $LSA(A)$ and continue with the following states l'_A until we either create an entirely new loop in $LSA(A)$ or generate l'_A already in $LSA(A)$ (we can stop generating l'_A as from now on, all l'_A are already in $LSA(A)$.

By executing the same steps for B, we get two LSAs, one for each finite automaton.

Running Example   For our finite automaton A_1, Figure 3.4 shows $LSA(A_1)$, currently prepared for an extraction of the length formulae for q_0. However, the initial state is irrelevant for the general $LSA(A_1)$, as it changes to the state we are currently computing length abstraction formulae for.


Figure 3.4: Summary lasso automaton $LSA(A_1)$.

3.1.3   Product Construction with Length Abstraction

The core of the product construction remains unchanged, but there are a few differences. The Algorithm 2 shows how we alternate the original product construction to optimize the algorithm with length abstraction.

We call W from line 3 a work set. It stores the potential product states prepared for processing, which we pick from W one by one[2].

The optimization process starts when we pick a product state p from W. Instead of immediately generating new successive product states p', we generate $lsa(q_1)$ and $lsa(q_2)$ to gain length formula. We test the satisfiability of this formula: $sat(\Phi^{LA}(p))$ where

$$\Phi^{LA}(p) : \alpha^{LA}(q_1) \wedge \alpha^{LA}(q_2) \text{ and}$$

$sat(\psi)$ is $True$ iff ψ is satisfiable (Φ **is** sat), $False$ otherwise. On line 9, we check whether $sat(\Phi^{LA}(p))$ holds and store a result as a boolean value to res. We are only interested in the satisfiability test result because we do not need any additional information from the computed formulae. Therefore, a simple boolean value is sufficient. $\Phi^{LA}(p)$ is passed to an SMT solver to solve its satisfiability. The α^{LA} compatibility check $\alpha^{LA}(q_1) \wedge \alpha^{LA}(q_2)$ **is** sat can be implemented in SMT solver as in Algorithm 3.   SMT solver returns sat when

---

[2]In spite of the fact more approaches are valid, we strongly recommend picking the last added product state from W (see line 7)—using depth-first search—as this allows us to quickly advance through the automaton and get to any final state faster—in case we just want to know whether automata have a non-empty intersection, this change will get us the answer most of the time in less steps. It works even better when implemented with a satisfiable state skipping optimization, explained in Section 3.1.4.

Input : NFA $A_1 = (Q_1, \Sigma, \delta_1, I_1, F_1)$, NFA $A_2 = (Q_2, \Sigma, \delta_2, I_2, F_2)$
Output: NFA $P = (A_1 \cap A_2) = (Q, \Sigma, \delta, I, F)$ with $L(P) = L(A_1) \cap L(A_2)$

1 $Q, \delta, F \leftarrow \varnothing$
2 $I \leftarrow I_1 \times I_2$
3 $W \leftarrow I$
4 $res \leftarrow False$
5 $solved \leftarrow \varnothing$
6 **while** $W \neq \varnothing$ **do**
7   **picklast** $[q_1, q_2]$ **from** W
8   **add** $[q_1, q_2]$ **to** $solved$
9 $res \leftarrow \alpha^{LA}(q_1) \wedge \alpha^{LA}(q_2)$ **is** sat
10   **if** $res = True$ **then**
11   **add** $[q_1, q_2]$ **to** Q
12   **if** $q_1 \in F_1$ **and** $q_2 \in F_2$ **then**
13   **add** $[q_1, q_2]$ **to** F
14   **forall** $a \in \Sigma$ **do**
15   **forall** $q_1' \in \delta_1(q_1, a), q_2' \in \delta_2(q_2, a)$ **do**
16   **if** $[q_1', q_2'] \notin solved$ **and** $[q_1', q_2'] \notin W$ **then**
17   **add** $[q_1', q_2']$ **to** W
18   **add** $[q_1', q_2']$ **to** $\delta([q_1, q_2], a)$

Algorithm 2: Product construction with length abstraction.

satisfiable (*res* is set to *True*) and *unsat* when unsatisfiable (*res* is set to *False*). If *unsat* is returned, length abstractions are incompatible. We have now pruned the generated state space by omitting the product state p.

1 `smtInit()`
2 `smtAdd(`$k \geq 0, m \geq 0$`)`
3 **for** $\varphi_{q_1} \in \alpha^{LA}(q_1)$ **do**
4   **for** $\varphi_{q_2} \in \alpha^{LA}(q_2)$ **do**
5   `smtPush()`
6   `smtAdd(`$\varphi_{q_1}.handle + \varphi_{q_1}.lasso * k = \varphi_{q_2}.handle + \varphi_{q_2}.lasso * m$`)`
7 $res \leftarrow$ `smtCheck()`
8   **if** $res = True$ **then**
9   **break**
10   `smtPop()`

Algorithm 3: Check compatibility of length abstractions with SMT solver.

If $sat(\Phi^{LA}(p))$, i.e., there will be an accepting run using p (see line 10), we add p to Q, possibly to F and generate p'.

A note of caution. It is important to understand that we are working only with possible word lengths and when we test the emptiness of the intersection of automata, we can resolve only such cases where words lengths are not accepted by both automata. When the test shows there could be some words of certain length accepted by both automata and for that reason by their intersection too—$sat(\Phi^{LA}(p))$—we cannot be sure there truly are any words accepted by both automata with their intersection non-empty, because there may be words of the suggested length, but it may be a different word for each automaton (which differ from one another in the containing symbols or their position in the word). For resolving such cases, we have to proceed with the classic algorithm steps to produce product states

according to their original transition symbols, not only by comparing the possible words lengths. With certainty, we can omit only the cases where $\neg sat(\Phi^{LA}(p))$.

Running Example We will continue with our running example. The second automaton we will be working with is a NFA $A_2 = (\{s_0, s_1, s_2, s_3\}, \{0, 1\}, \delta_2, \{s_0\}, \{s_3\})$ where δ_2 is depicted in Figure 3.5.


Figure 3.5: Non-deterministic finite automaton A_2.

In Figure 3.6, there is $LSA(A_2)$, which we will be using together with $LSA(A_1)$ shown in Figure 3.4 for product construction of A_1 and A_2.


Figure 3.6: Lasso automaton $LSA(A_2)$ for A_2.

When we start the algorithm, we get the following length abstraction formulae for $p = [q_0, s_0]$. From $LSA(A_1)$ for q_0 (q_0 is the new initial state of $LSA(A_1)$), we get an existential formula representing length abstraction $\alpha^{LA}(q_0)$[3]. From $LSA(A_2)$ for s_0 (s_0 is the new initial state of $LSA(A_2)$), we get a formula for length abstraction $\alpha^{LA}(s_0)$[4].

$$\alpha^{LA}(q_0) : \exists k(|w| = 2 \lor |w| = 4 + 2 \cdot k)$$
$$\alpha^{LA}(s_0) : \exists m(|w| = 2 + 1 \cdot m)$$

When we compare $\alpha^{LA}(q_0)$ and $\alpha^{LA}(s_0)$, we get:

$$\alpha^{LA}(q_0) \land \alpha^{LA}(s_0) : \exists k(|w| = 2 \lor |w| = 4 + 2 \cdot k) \land \exists m(|w| = 2 + 1 \cdot m)$$

or in a simplified notation:

$$\alpha^{LA}(q_0) \land \alpha^{LA}(s_0) : \exists k \exists m(2 \lor 4 + 2 \cdot k = 2 + 1 \cdot m).$$

To solve satisfiability of $\Phi^{LA}([q_0, s_0])$, we try to find values of k and m such that $|w|$ in both formulae are equal (some expressions on the left and on the right side of the equation are equal).

---

[3]This formula consists of two independent disjuncts φ_1 and φ_2 describing there are more possible lengths for accepted words from the same initial state.

[4]We are using variable m here instead of k to emphasize variables from different formulae are not dependent on each other—they belong to different LSAs.

In Figure 3.7, we can see the product of A_1 and A_2 being constructed using length abstraction. Red states represent product states whose formulae are resolved as unsatisfiable and therefore the algorithm omits any successive product states—dashed states (such as q_4s_2 or q_3s_2) which are generated in the unoptimized product construction. The green state represents final states in both automata. Here, we have found a solution accepted by both A_1 and A_2. If we desire to resolve only the emptiness problem, we can stop the execution of the algorithm here as we have found one final state—automata have non-empty intersection. The blue state is a normal product state whose significance will be explained in section 3.1.4.


Figure 3.7: Constructed product with depiction of length abstraction optimization.

As you can notice in Figure 3.8, the product generated by our algorithm has only 4 product states in comparison to 9 product states generated by the unoptimized product construction.


Figure 3.8: Final product minimized by length abstraction.

3.1.4   Optimization with Skipping Satisfiable States

When we take new p from W and $sat(\Phi^{LA}(p))$, it is time to add all the possible successive product states p' to W. When p generates only a single p' and $p \notin F_P$ (final states are obviously in the product) or satisfiable length was not zero[5], we can say with certainty that $sat(\Phi^{LA}(p'))$ as there is only a single branch in the automaton leading from p to a final state (through p'). p' is *skippable*, iff there exists $p \notin F_P$ or with satisfiable length not zero for which $sat(\Phi^{LA}(p))$ and whose only successor is p', we add p' to W with the information of being skippable. If p' is already in W, we append the information to p' in W.

We skip checking for $sat(\Phi^{LA}(p'))$ when we pick p' from W. We can immediately check for final states and generate the successive product states. This optimization saves

---

[5]If the satisfiable length is zero, we are in a final product state and p generated from final state might not lead to any final state, but the length abstractions for final state are compatible, of course.

us generating the length abstraction formulae for p' and testing the formulae in the SMT solver for their satisfiability.

If automata have long lines (with non-splitting branches), this will prove extremely useful, because only a few proper iterations with formulae computing and running SMT solver computation will be executed. The application of skipping satisfiable states is depicted in Algorithm 4. The line 9 from Algorithm 2 is substituted with the contents of Algorithm 4.

1 **if** skippable($[q_1, q_2]$) **then**
2 $\quad res \leftarrow True$
3 **else**
4 $\quad res \leftarrow \alpha^{LA}(q_1) \wedge \alpha^{LA}(q_2)$ **is** sat

Algorithm 4: Substitution of line 9 in Algorithm 2 with skipping satisfiable states.

The only change is a test for every checked p, which decides whether p can be skipped. You can see that we proceed with the satisfiability check in SMT solver only for p which are generated from the product states with multiple transitions generating p and at least one more product state (in general at least two new potential product states). If only one p was generated earlier from a product state with satisfiable formulae, we skip the check for $sat(\Phi^{LA}(p))$ and continue to generating its successive states immediately.

You can notice there is one skippable state in the former example, which had to be evaluated and tested for satisfiability earlier. The blue state in Figure 3.7 is such a skippable state. In our case for state $q_5 s_2$, when only one new state is generated from state $q_4 s_3$ while this state is resolved as satisfiable (with not zero length—otherwise, if $q_5 s_2$ did not lead back to $q_4 s_3$, $q_5 s_2$ would be skippable even though it would not lead to any final state), newly generated product state has to be satisfiable as well, because the check for $q_4 s_3$ already considered the state $q_5 s_2$ as its only way to any final state with not zero length.

When we have a series of such states, we can highly optimize generating the whole branch with only one initial check for satisfiability. In real world examples, there are often automata with long branches splitting into multiple branches only occasionally. We will check for satisfiability for all the initial states of each new branch and then either omit the entire branch (if *unsat* is returned) or skip checking satisfiability in the entire branch (if *sat* is returned).

3.1.5 Resolving Length Abstraction Satisfiability without SMT Solver

Evaluating satisfiability of length abstractions formulae in SMT solver is expensive. We try to replace SMT solver with a specialized structure which transforms the problem of solving satisfiability of length abstraction formulae to evaluating satisfiability of a linear congruence equations.

Length abstraction formulae have the same, simple structure. Length abstraction can be implemented as a set of length formulae represented as a two-tuple of handle length and lasso length. Therefore, we can easily compare such sets in order to resolve satisfiability of length abstraction formulae without SMT solver. $\Phi^{LA}(p)$ forms a set of linear congruence equations, which can be resolved just by utilizing basic mathematical operations and properties of linear congruences.

The Algorithm 5 shows how to determine $sat(\Phi^{LA}(p))$ from line 9 using linear congruences.

We execute the following steps for each equation

$$\varphi_{q_1}.handle + \varphi_{q_1}.lasso \cdot k = \varphi_{q_2}.handle + \varphi_{q_2}.lasso \cdot m.$$

```
1  for φ_{q_1} ∈ α^{LA}(q_1) do
2  │    for φ_{q_2} ∈ α^{LA}(q_2) do
3  │    │    if φ_{q_1}.handle = φ_{q_2}.handle then
4  │    │    │    res ← True
5  │    │    else if φ_{q_1}.handle > φ_{q_2}.handle then
6  │    │    │    res ← solveForOneHandleLonger(φ_{q_1}, φ_{q_2})
7  │    │    else
8  │    │    │    res ← solveForOneHandleLonger(φ_{q_2}, φ_{q_1})

9  res ← False
```
Algorithm 5: Check satisfiability using length abstraction algorithm without SMT solver.

If the handle lengths of φ_{q_1} and φ_{q_2} are equal, there are words of the same length accepted by both $\alpha^{LA}(q_1)$ and $\alpha^{LA}(q_2)$ (they are mutually compatible) without stepping into the loops of $LSA(A_1)$ and $LSA(A_2)$. Otherwise, handle lengths differ, and we must consider lengths of loops in our determination of compatibility of $\alpha^{LA}(q_1)$ and $\alpha^{LA}(q_2)$.

We now have to determine $sat(\varphi_{q_1} \wedge \varphi_{q_2})$ for abstractions with one handle longer. This is solved by the function `solveForOneHandleLonger` in Algorithm 6.

First, to simplify the equation (line 2), we move handle lengths from the side of the equation with the shorter handle φ_s to the side with the longer handle φ_l to solve:

$$\varphi_l.handle + \varphi_l.lasso \cdot k = \varphi_s.lasso \cdot m \tag{3.1}$$

which represents the number of loops a word must make in φ_s to be accepted by φ_l (as if with shorter $\varphi_l.handle$).

If both φ_l and φ_s have no loops (line 4), $\varphi_{q_1} \wedge \varphi_{q_2}$ is unsatisfiable because the handles differ. Else, if only φ_s has no loop (line 6), every word accepted by φ_s is shorter than words accepted by φ_l and the formulae cannot be satisfiable.

Else, if only φ_l has no loop (line 8), we can try to manually iterate over loops in φ_s to see whether the difference of word lengths between handles can be equalized by looping in $\varphi_s.lasso$.

Otherwise, both φ_l and φ_s have loops (line 16). We can apply linear congruence properties to the equation 3.1 to determine whether the formulae are satisfiable. The equation 3.1 says that if formulae are satisfiable, the left side of the equation is divisible by some multiple of $\varphi_s.lasso$. We can rewrite that in a linear congruence equation as follows:

$$\varphi_l.handle + \varphi_l.lasso \cdot k \equiv 0 \pmod{\varphi_s.lasso} \tag{3.2}$$
$$\varphi_l.lasso \cdot k \equiv -\varphi_l.handle \pmod{\varphi_s.lasso} \tag{3.3}$$

which is the same as solving a linear Diophantine equation

$$\varphi_l.lasso \cdot k - \varphi_s.lasso \cdot m = -\varphi_l.handle. \tag{3.4}$$

Properties of multiplicative inverse [9, 19], based on Bézout's identity [9], say that iff $\varphi_l.lasso$ and $\varphi_s.lasso$ are relatively prime (coprime)—the greatest common divisor (GCD) of $\varphi_l.lasso$ and $\varphi_s.lasso$ is equal to 1—there exists a multiplicative inverse for $\varphi_l.lasso$ in modulo $\varphi_s.lasso$ which ensures that linear congruence 3.3 is always solvable for some $\varphi_l.lasso$ in modulo $\varphi_s.lasso$[6]. Therefore, the formulae are satisfiable.

Otherwise, $\varphi_l.lasso$ and $\varphi_s.lasso$ are not coprime (GCD is different from 1) and by properties of linear Diophantine equations [19], iff GCD precisely divides y without a remainder

---

[6]We can get the precise solution by multiplying both sides of the equation with the multiplicative inverse.

```
1  Function solveForOneHandleLonger(φ_l, φ_s):
       Data: Input length abstraction formulae of potential product state.
       φ_l: Length abstraction formula with the longer handle,
       φ_s: Length abstraction formula with the shorter handle.
       Result: bool: True if satisfiable, False otherwise.
2      φ_l.handle ← φ_l.handle − φ_s.handle
3      φ_s.handle ← 0
4      if φ_l.lasso = 0 and φ_s.lasso = 0 then
5      │   return False
6      else if φ_s.lasso = 0 then
7      │   return False
8      else if φ_l.lasso = 0 then
9      │   it ← 0 // Current length resembling the iteration of the shorter lasso loop.
10     │   while it ≤ φ_l.handle do
11     │   │   if it = φ_l.handle then
12     │   │   │   return True
13     │   │   else
14     │   │   │   it ← it + φ_s.lasso
15     │   return False
16     else
17     │   gcd ← getGCD(φ_l.lasso, φ_s.lasso)
18     │   if gcd = 1 then
19     │   │   return True
20     │   else
21     │   │   y ← −φ_l.handle
22     │   │   while y < gcd do
23     │   │   │   y ← y + φ_s.lasso
24     │   │   if (y mod gcd) = 0 then
25     │   │   │   return True
26     │   │   else
27     │   │   │   return False
```

Algorithm 6: Solve satisfiability of length abstraction formulae for one handle longer.

where y is the right side of the linear congruence 3.3 or its any congruent equivalent, there exist solutions to the linear congruence[7]. Otherwise, there are no solutions.

3.2   Parikh Image Abstraction of State Languages

Length abstraction is a simple and fast optimization, but can be too coarse to detect non-terminating states in some cases. In this section, we present an abstraction of state languages with Parikh images, α^{PI}, which aims to replace length abstraction to make the abstraction more precise to prune larger quantities of product state space.

Parikh images provide more information about the finite automata than simple length abstraction. While length abstraction considers only accepted word lengths without knowing which transition symbols are actually in the transitions, Parikh image abstracts the state language to numbers of occurrences of specific transition symbols in words regardless of their position in said words. Thus, Parikh image abstraction allows us to more precisely determine whether the product state has non-empty language. However, Parikh image

---

[7]We can apply extended Euclidean algorithm to find the precise values for the Diophantine equation 3.4.

computation itself is expensive. The question is, whether the added computation time compensates for more precise product generation with higher state pruning capabilities.

We will introduce an algorithm for Parikh image abstraction α^{PI} applied on each product state $p = [q_1, q_2]$ to decide the compatibility of $\alpha^{PI}(q_1)$ and $\alpha^{PI}(q_2)$.

3.2.1 Parikh Image

We derive our Parikh image construction from the Parikh's theorem [25] described in [13], creating a semi-linear Parikh image formulae for the given regular language as a set of Parikh images for each word in the language. However, our usage of Parikh image of some regular language (and therefore of the corresponding finite automaton recognizing such regular language) is restricted to determining the compatibility of Parikh image state language abstractions. Therefore, we only test for satisfiability of Parikh image formulae describing $\alpha^{PI}(q)$. We use SMT solver to resolve the satisfiability of Parikh image formulae of the current potential product state.

Given an NFA $A = (Q, \Sigma, \Delta, I, F)$, Parikh image formula φ (as described in [24] for solving string constraints) consists of several constraints in conjunctive normal form. φ describes runs of A. Each satisfiable assignment defines properties of the run. φ consists of the following conjuncts:

1. Foremost, we define a variable u_q for each state $q \in Q$. u_q defines how many times we enter q and exit q by specifying the difference between the number of entries and exits. We construct equations with u_q for a run as follows:

   - $u_q = 1$ for $q \in I$,
   - $u_q \in \{0, -1\}$ for $q \in F$ and
   - $u_q = 0$ for $q \in Q \smallsetminus (I \cup F)$.

2. Second, we define a variable y_t for each transition $t \in \Delta$ such that $y_t \geq 0$ describing how many times is t used in the run.

3. We can now present an equation introducing a connection between u_q and y_t to evaluate the difference between the number of entries and exits for each $q \in Q$ as follows:
$$u_q + \sum_{t \in \Delta_q^+} y_t - \sum_{t \in \Delta_q^-} y_t = 0.$$

   where Δ_q^+ is a set of ingoing transitions $\Delta_q^+ = \{(q', a, q) \in \Delta\}$ and Δ_q^- is a set of outgoing transitions $\Delta_q^- = \{(q, a, q') \in \Delta\}$ from the given state q.

4. Furthermore, we need to make sure that the states used in runs described by the satisfying assignments are connected and start in the initial state. Variable z_q for each $q \in Q$ is introduced. z_q represents the length of any path from I to q in a spanning tree of the subgraph with $y_t \geq 0$. If $z_q = 0$, there is no path from I to q and the state q is not used in the run. $z_q > 0$ means there is a path from I to q and q is used in the run.

   If $q \in I$, we add a constraint $z_q = 1 \wedge y_t \geq 0$. Otherwise,
$$\left(z_q = 0 \wedge \bigwedge_{t \in \Delta_q^+} y_t = 0\right) \vee \bigvee_{t \in \Delta_q^+} (y_t \geq 0 \wedge z_{q'} \geq 0 \wedge z_q = z_{q'} + 1).$$

   If the distance z_q is 0, q is not in the run.

5. Last but not least, we declare the only free variable $\#_a$ for each transition symbol $a \in \Sigma$. $\#_a$ describes the number of occurrences of a in accepted words regardless of their position in the words (the number of a in the run). $\#_a$ is the only variable common to different Parikh image abstractions when we test their compatibility. The constraint $\#_a = \sum_{t=(q,a,q') \in \Delta} y_t$ ensures $\#_a$ is consistent with the number of used t with a.

We gain an existentially quantified formula φ in Presburger arithmetic describing language abstracting α^{PI} for A with free variables $\#_a$:

$$\alpha^{PI} : \exists u_{q_1}, \ldots, u_{q_n}, z_{q_1}, \ldots, z_{q_n}, y_{t_1}, \ldots, y_{t_m}(\varphi)$$

where $n = |Q|$ is the number of states and $m = |\Delta|$ is the number of transitions in the finite automaton.

Notice that α^{PI} is an existential formula, which is great for SMT solving where computing with universal quantifiers can take a long time. SMT solver are specialized on efficient solving of existential or quantifier-free formulae.

As for length abstraction for product state $p = [q_1, q_2]$, we decide compatibility of Parikh image formulae $\alpha^{PI}(q_1)$ and $\alpha^{PI}(q_2)$ as follows: $sat(\Phi^{PI}(p))$ such that

$$\Phi^{PI}(p) : \alpha^{PI}(q_1) \wedge \alpha^{PI}(q_2)[8].$$

3.2.2 Product Construction with Parikh Image Abstraction

We introduce the unoptimized product construction using Parikh image abstraction. The algorithm is analogous to the product construction optimized by length abstraction from Algorithm 2. The difference is that we now compute Parikh image formulae and determine their satisfiability instead of generating lasso automata and determining satisfiability of length abstraction formulae.

We use Parikh image formulae to determine whether p is to be added to the product P. As for length abstraction, we test whether Parikh image abstractions are compatible (a conjunction of Parikh image formulae is satisfiable). Therefore, instead of length abstraction on line 9 in Algorithm 2, we compute Parikh image abstractions: line 9 is replaced with

$$res \leftarrow \alpha^{PI}(q_1) \wedge \alpha^{PI}(q_2) \text{ is } sat$$

We can see our proposed algorithm using Parikh image computation to optimize product construction in the Algorithm 7. Parikh image formulae are computed on line 9 and their satisfiability is determined.

3.2.3 Reduced Parikh Image

The presented Parikh image would work well regarding its pruning capabilities. However, the described Parikh image computation requires extensive resources and computation time and we need Parikh images computed only for determining the emptiness of the intersection. Given that most of the computation time is spent by the evaluation of Parikh image

---

[8]In reference implementation, we replace existential formulae with quantifier-free formulae with renamed variables without existential quantifiers.

Input : NFA $A_1 = (Q_1, \Sigma, \delta_1, I_1, F_1)$, NFA $A_2 = (Q_2, \Sigma, \delta_2, I_2, F_2)$
Output: NFA $P = (A_1 \cap A_2) = (Q, \Sigma, \delta, I, F)$ with $L(P) = L(A_1) \cap L(A_2)$

1   $Q, \delta, F \leftarrow \varnothing$
2   $I \leftarrow I_1 \times I_2$
3   $W \leftarrow I$
4   $res \leftarrow False$
5   $solved \leftarrow \varnothing$
6   **while** $W \neq \varnothing$ **do**
7   **picklast** $[q_1, q_2]$ **from** W
8   **add** $[q_1, q_2]$ **to** $solved$
9       $res \leftarrow \alpha^{PI}(q_1) \wedge \alpha^{PI}(q_2)$ **is** sat
10   **if** $res = True$ **then**
11   **add** $[q_1, q_2]$ **to** Q
12   **if** $q_1 \in F_1$ **and** $q_2 \in F_2$ **then**
13   **add** $[q_1, q_2]$ **to** F
14   **forall** $a \in \Sigma$ **do**
15   **forall** $q_1' \in \delta_1(q_1, a), q_2' \in \delta_2(q_2, a)$ **do**
16   **if** $[q_1', q_2'] \notin solved$ **and** $[q_1', q_2'] \notin W$ **then**
17   **add** $[q_1', q_2']$ **to** W
18   **add** $[q_1', q_2']$ **to** $\delta([q_1, q_2], a)$

Algorithm 7: Product construction with Parikh image abstraction.

conjuncts in SMT solver, we want to minimize the number of Parikh image conjuncts SMT solver needs to evaluate for each φ.

Consequently, we infer our reduced Parikh image from the shown Parikh image to further optimize Parikh image computation. We modify several conjuncts in Parikh image formula and unify initial states and accepting states to simplify the formula and reduce its complexity.

Due to how we have reduced our Parikh image, we work only with finite automata with a single initial state and a single accepting state. However, we can easily convert any finite automaton into the required format with adding two new states: one for a new initial state from which one can transition to all previous initial states and one for a new accepting state to which lead all previous accepting states. The previous initial and accepting states are changed to common automata states.

Our reduced Parikh image consists of the following conjuncts:

1. We use the conjuncts 1, except now we restrict u_q for each final state to have only the value -1, i.e.:

$$u_q = -1 \text{ for each state } q \in F.$$

   We can perform this reduction, because we know for sure that by unifying final states of the automaton into one abstract final state, there will be exactly only one final state where all words accepted by the automaton end, but none passes through this state earlier.

2. The conjuncts 2 and 3 remain unchanged, the same holds for conjuncts 5.

3. However, we completely omit the conjuncts for z_q. The reason is that, as we have found out, the difference in pruning capabilities of Parikh image with or without the conjuncts 4 on our benchmark automata is insignificant in comparison to the computation time spared by removing these conjuncts.

The reason conjuncts 4 are so computationally costly is that they are complex for even simple automata. Even then, if we want to keep them, we can include these conjuncts, but, we can reduce their complexity by not having to compute z_q lengths for initial and final states.

The constraint for when q is an initial state ($z_q = 1 \wedge y_t \geq 0$) remains unchanged as a starting length for other states. However, for every other state, we remove the possibility of $y_t = 0$ and $z_{q'} = 0$ in the second half of the conjuncts (as the option cannot occur with unified initial and final states). The conjuncts look like this:

$$(z_q = 0 \wedge \bigwedge_{t \in \Delta_q^+} y_t = 0) \vee \bigvee_{t \in \Delta_q^+} (y_t > 0 \wedge z_{q'} > 0 \wedge z_q = z_{q'} + 1).$$

Skippable States Optimization

Same as for the length abstraction, we can make use of skipping satisfiable product states optimization. When $sat(\Phi^{PI}(p))$ for some potential product state $p = [q_1, q_2]$ and p generates only one consecutive potential product state $p' = [q'_1, q'_2]$ such that $p \xrightarrow{a} p'$ where $a \in \Sigma$, we can skip computing Parikh images for p' as we know for sure $sat(\Phi^{PI}(p'))$ in order to get a satisfiable result for Parikh image for p. We can add this functionality to our previous algorithm by replacing line 9 with the content of Algorithm 8.

1 **if** skippable($[q_1, q_2]$) **then**
2 $res \leftarrow True$
3 **else**
4 $res \leftarrow \alpha^{PI}(q_1) \wedge \alpha^{PI}(q_2)$ **is** sat

Algorithm 8: Parikh image computation with skippable states optimization.

3.2.4   Optimization with Incremental SMT Solving

Parikh image formulae are large and SMT solving is expensive. We have to recompute Parikh image formulae for every potential product state. However, formulae generated for different product states in one intersection problem are very similar. Large parts of Parikh image formulae do not change between the product states at all.

We try to use SMT solver with incremental SMT solving to reuse parts of the previous computation in the next one. We can specify parts of the Parikh image formulae in the SMT solver once, without passing them to the solver for each product state. Further, once a formula have been computed, the solver can use its cache to reuse parts of the computation[9]. In this section, We explain how we use incremental SMT solving for Parikh images in product construction to compute similar, consecutive Parikh image formulae faster.

Notice that some conjuncts of Parikh image remain unchanged for the whole automaton, i.e., for every product state. Only some conjuncts which work with initial states (conjuncts 1 and 4) have to be rewritten, because the only difference between states in two different product states are the different initial states.

Assume finite automata A_1 and A_2 (whose intersection we generate) and a product state $p = [q, s]$ where $q \in Q_{A_1}, s \in Q_{A_2}$. The changes of conjuncts in φ_{A_1} and φ_{A_2} are caused by moving (setting) the states in both A_1 and A_2 corresponding to p as new initial states

---

[9]Consequently, the computation of the first Parikh image takes longer than for the next states.

$I_{A_1} = \{q\}$ and $I_{A_2} = \{s\}$ as we proceed further into the automata with product construction. We start with the abstract initial states (one for each original automata, $I_{A_1} = \{q'_0\}$ and $I_{A_2} = \{s'_0\}$).

First, we compute $\Phi^{PI}(p_0)$ such that $p_0 = [q'_0, s'_0]$. Iff $sat(\Phi^{PI}(p_0))$, we generate new potential product states (e.g., $p_1 = [q_1, s_1]$ and $p_2 = [q_1, s_2]$). Now we need to check whether to include p_1 and p_2 to the generated product, i.e., check that $sat(\Phi^{PI}(p_1))$ and $sat(\Phi^{PI}(p_2))$, respectively. Taking p_1, we set new initial states $I_{A_1} = \{q_1\}, I_{A_2} = \{s_1\}$. Similarly, for p_2, we would set $I_{A_1} = \{q_1\}, I_{A_2} = \{s_2\}$.

We now need to change every mention of initial states in φ_{A_1} and φ_{A_2} because the initial states are different from those we used at the start (q'_0 and s'_0) and for which we already computed $\Phi^{PI}(p_0)$. We now introduce an optimization of Parikh image computation which precomputes unchanged conjuncts only once and recomputes only conjuncts mentioning initial states.

Persistent and State Specific Clauses

To present optimization with incremental SMT solving, we split $\alpha^{PI}(q)$ conjuncts into two groups: persistent clause and state specific clause.

Persistent clause represents Parikh image conjuncts which can be precomputed once for all states in the finite automaton and used throughout the whole product construction. Persistent clause consists of unchanged conjuncts of reduced Parikh image described in 3.2.1: conjuncts 2, conjuncts 3 and conjuncts 5.

The state specific clause consists of conjuncts which change with every product state p, and as such have to be constructed and recomputed for every satisfiability test. The process of recomputing state specific clauses is the most expensive part of the product construction algorithm using Parikh images. Therefore, our goal is to minimize the number of conjuncts in a state specific clause as much as possible. The state specific clause consists of conjuncts 1 in reduced Parikh image as they directly change according to initial states and, optionally, if we want to include z_q conjuncts, conjuncts 3. We would need to recompute z_q conjuncts for each potential product state too because the conjuncts compute with initial states.

It is worth to note that the conjuncts 3 in reduced Parikh image manipulate with initial states but the structure of the conjuncts could be reversed to compute connectedness of the automaton in *reversed* order, from the accepting states to the initial states. In that case, the conjuncts could be reconstructed as a part of the persistent clause dependent on accepting states which remain unchanged (the abstract accepting state) for the entire time. This additional optimization might be worth inspecting. Because the inclusion of conjuncts 3 does not generate smaller state spaces with our benchmark automata, we did not investigate further yet.

Algorithm for Incremental SMT solving Using Parikh Image

To implement incremental SMT solving to our current Parikh image computation shown in Algorithm 7, we need to make the following adjustments.

We need to precompute persistent clauses once for both A_1 and A_2. We insert a new line to our algorithm between lines 5 and 6. The new line contains a call to a function `addPersistentClauses()` which precomputes persistent clauses for both A_1 and A_2. Note that the function is called only once, before we enter the *while* loop for iterating over potential product states.

We compute state specific clauses as normal when we ask whether $sat(\Phi^{PI}(p))$ when we are checking compatibility of both α^{PI} on line 9. However, we push the previously precomputed state persistent clauses to the SMT solver stack. This preserves them when the current state specific clauses are dropped after $sat(\Phi^{PI}(p))$ is resolved. For a pseudocode of the replacement of line 9, see Algorithm 9.

1 `smtPush()`
2 $res \leftarrow \alpha^{PI}(q_1) \wedge \alpha^{PI}(q_2)$ **is** sat
3 `smtPop()`
Algorithm 9: Add state specific clauses to SMT solver for incremental SMT solving optimization.

The line 2 computes Parikh image formulae and determines their satisfiability, as explained in Section 3.2.3.

3.2.5   Optimization with SMT Solver Timeout

In the case of Parikh images computed with SMT solver, it is easier to determine $\neg sat(\Phi^{PI}(p))$ than $sat(\Phi^{PI}(p))$. Based on our experiments, we use timeout functionalities of SMT solver to speed up the process of resolving satisfiability of potential product states.

We define a maximal amount of time SMT solver can compute $sat(\Phi^{PI}(p))$ for a single product state p to resolve its satisfiability. If SMT solver resolves $sat(\Phi^{PI}(p))$ before the time runs out, we proceed as normal. However, if the time runs out, the result of the satisfiability test is unknown and we must presume $\Phi^{PI}(p)$ could be satisfiable: we must set res to $True$.

This approach resolves $sat(\Phi^{PI}(p))$ of an over-abstraction described previously. We prune such potential product states that $sat(\Phi^{PI}(p))$ can be resolved quickly (within the defined timeout) while allowing the inclusion of some potential product states which are in fact unnecessary to the generated product. Nevertheless, we find pruning capabilities of this optimization satisfactory and the computation time decreases noticeably.

The timeout is chosen empirically. One has to experiment with their finite automata. The ideal timeout can vary for different benchmarks. One timeout is usually successfully usable for operations on similar finite automata. The timeout is directly proportional to a precision of Parikh image abstraction and reversely proportional to the scale of Parikh image over-abstraction.

3.3   Combination of State Language Abstractions

Length abstraction is fast but coarse; Parikh image abstraction is precise but expensive. We can combine both abstractions to take advantage of respective strengths of our abstractions. In this section, we present an algorithm which introduces a modification to evaluation of compatibility of state language abstractions. We use both length abstraction and Parikh image computation to determine satisfiability of state abstraction to optimize product construction. The pruning capabilities remain the same as if we computed Parikh image alone, or even better in cases where Parikh image computation times out.

The Algorithm 10 shows how we apply our modifications on a single evaluation of compatibility of abstractions.

First, we test whether α^{LA} alone can prune the generated product state space by omitting the current potential product state $[q_1, q_2]$ if $\neg sat(\Phi^{LA}([q_1, q_2]))$. If length abstraction

```
1  if α^{LA}(q_1) ∧ α^{LA}(q_2) is unsat then
2  │   res ← False
3  else
4  │   res ← α^{PI}(q_1) ∧ α^{PI}(q_2) is sat
5  │   if res = Unknown then
6  │   │   res ← True
```

Algorithm 10: Implementation of checking compatibility of state abstractions using both length abstraction and Parikh image computation optimizations.

succeeds in omitting $[q_1, q_2]$ from the product, we do not need to compute Parikh images for $[q_1, q_2]$ and can continue with the product construction as if $\neg\Phi^{PI}([q_1, q_2])$. Otherwise, we continue with Parikh image computation for $[q_1, q_2]$ (resolving satisfiability of its formulae as in the basic Parikh image algorithm from Algorithm 7).

3.4 Abstraction of State Languages with Mintermization

In this section, we introduce a method of optimizing operations on finite automata using minterms [23]. Minterm computation abstracts the state language of automata differently than what we have explored so far, allowing us to follow a diverse set of characteristics about the state language. We can afterwards make use of computed minterms for the automata with other optimization methods introduced in this paper, as well as another optimization approaches.

Foremost, we give an algorithm for minterm computation adapted from [11], further defined and expanded in [22] for simulation algorithms for symbolic automata and now optimized for product construction to compute minterms for the non-empty multiset of input finite automata $A = \{A_1, A_2, \ldots, A_n\}$ where n equals the number of finite automata. Gained minterms abstract automata state language in such a way we do not lose any information about the original automata (minterms are not an over-approximation of the original automata), but might create a more concise finite automata which will be easier to work with in our other abstractions and may significantly decrease the computation time required for optimizations such as Parikh image computation.

The general idea is to get sets of transition symbols between two states for all our considered finite automata. Compute minterms from these sets and substitute transition symbols between two states in our automata with corresponding minterms created from these transition symbols.

For now, let us explain what minterms are and how you can generate them.

Definition 3.4.1 (Minterms)
Given an NFA $A = (Q, \Sigma, \delta, I, F)$, let $\Phi = \{\varphi_1, \varphi_2, \ldots, \varphi_n\}$ be a finite set of non-empty finite sets of transition symbols $\varphi_i = \{a \mid a \in \Sigma \wedge q \xrightarrow{a} q'\}$ for $1 \leq i \leq n$ where $q, q' \in Q$, n equals the number of state pairs (q, q') such that $q \xrightarrow{a} q'$ where $q' \in \delta(q, a)$.

We call φ_i a *transition set* for the given pair of automaton states q, q'. We denote Ψ or Minterms(Φ) as a set of all minterms ψ for A such that

$$\Psi = Minterms(\Phi) = \left\{ \psi = \bigcap_{1 \leq i \leq n} \psi_i \;\middle|\; \forall i \in \{1, \ldots, n\}\big((\psi_i \in \{\varphi, Q \smallsetminus \varphi\}) \wedge \psi \neq \varnothing\big) \right\}.$$

Minterms are computed once, at the beginning of the optimization process for all considered finite automata. We generate so called *minterm tree* with nodes as intersection between

25

sets of transition symbols in case the intersection is non-empty. Each node can have up to two children, representing intersection with the next transition set and its complement, respectively.

When such minterms for the given automaton are computed, we can abstract the state language of the automaton by replacing transitions from the state by their corresponding minterms. We say minterm ψ is created from the set of transition symbols $\varphi \in \Phi$ if φ is used in the intersection defining ψ in its direct form, not as a complement $Q \smallsetminus \varphi$.

Notice that we can compute minterms over multiple NFAs, which allows us to use minterms state language abstraction for optimization of operations on those automata.

Given finite automata $A_1 = (\{s_0, s_1, s_2, s_3\}, \Sigma, \delta_1, \{s_0\}, \{s_3\})$ and $A_2 = (\{q_0, q_1, q_2\}, \Sigma, \delta_2, \{q_1\}, \{q_0\})$ over alphabet $\Sigma = \{a, b, c, d\}$ with δ_1 and δ_2 according to Figure 3.9 and Figure 3.10, respectively, the Figure 3.14 depicts how we could mark each transition set in our automata to be used in mintermization process. For example, a transition set φ_1 could be a set of transition symbols from state s_0 to s_1: $\varphi_i = \{a, b, d\}$. Similarly, we mark the remaining transition sets. Now, we can proceed to execute mintermization operations.

Figure 3.9: Finite automaton A_1 with transitions δ_1.

Figure 3.10: Finite automaton A_2 with transitions δ_2.

Figure 3.11: Finite automata A_1 and A_2 used as example automata for mintermization.


Figure 3.12: Finite automaton A_1 with transition sets φ_i.

Figure 3.13: Finite automaton A_2 with transition sets φ_i.

Figure 3.14: Finite automata A_1 and A_2 with marked transition sets used in mintermization.

Computation of minterms for A_1 and A_2 is illustrated in Figure 3.15 in a diagram.
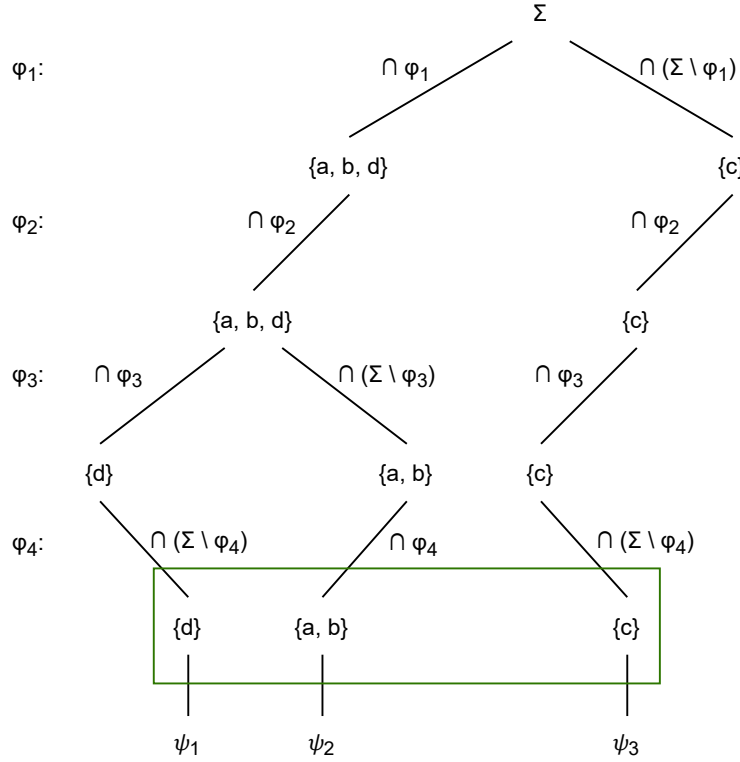
Figure 3.15: Mintermization process executed on example finite automata A_1 and A_2. We start with the whole alphabet and make our way down through all mintermization sets φ_i, where $1 \leq i \leq n$. For each mintermization set, we compute the intersection of the preceding set with the current mintermization set φ_i. The results are shown in the diagram as the nodes of the tree. When operations on all mintermization sets were executed, the leaves of the tree (indicated by the green square) represent the final minterms for the given mintermization sets Φ over the given alphabet Σ. We denote each minterm ψ_i, where $1 \leq i \leq |\Psi|$ where $|\Psi|$ represents the total number of generated minterms.

We start with the whole alphabet of both automata[10] at the top of the minterm tree to be generated. Afterwards, we iterate over transition sets. For each transition set φ_i, we compute the intersection of the current minterm tree leaves with:

- the current transition set φ_i and store the result as a left node of this particular tree node,

- the complement of the current transition set $Q \smallsetminus \varphi_i$ and store the result as a right tree node of this particular tree node.

If the intersection is empty, we omit creating the corresponding child node entirely. In the end, we are left with a complete minterm tree for the given set of transition sets Φ representing the specified finite automata.

---

[10]If the automata had non-equal alphabets, we would start with their intersection: $\Sigma = \Sigma_1 \cap \Sigma_2$. This is an optimization specific to product construction: If some transition symbols are not used by every finite automaton, we can safely omit such symbols as they are definitely not present in the intersection of these automata.

The acquired minterms are:

$$\Psi = Minterms(\Phi) = \big\{\{d\}, \{a, b\}, \{c\}\big\} = \{\psi_1, \psi_2, \psi_3\}.$$

We can now substitute the former transition sets φ_i for finite automata with the appropriate minterms $\psi_j, 1 \le j \le |\Psi|$ which were created from the specific transition sets $\varphi_i \in \Phi$ such that φ_i is used in its direct form (not as a complement) in the process of computing ψ_j (optimized for product construction). The gained automata can be seen in Figure 3.18.

Figure 3.16: Finite automaton A_1 with transitions substituted by corresponding minterms $\psi_i \in \Psi$ created from these transition sets.

Figure 3.17: Finite automaton A_2 with transitions substituted by corresponding minterms $\psi_i \in \Psi$ created from these transition sets.

Figure 3.18: Finite automata A_1 and A_2 with substituted transitions with minterms in the process of mintermization.

As we can see, we are able to get rid of some transition symbols and reduce the alphabet as well as the number of transitions in finite automata. Considering we have minterms over alphabet of A, we know that the intersection of two minterms has to be an empty set and that $\forall \psi \in \Psi (\psi \subseteq \varphi, \varphi \in \Phi)$ if ψ is created from φ. Important improvement of using minterms in product construction is the fact that $|\Psi| \le |\Sigma|$ instead of at most $2^{|\Sigma|}$ as is the case for minterms over general predicates for general operation (e.g., [22]). We make use of these points further.

We can use the method of minterm computation with length or Parikh image abstractions of state languages. We choose this approach in order to improve pruning capabilities of length abstraction for some finite automata or speed up demanding Parikh image computation, especially for automata with multitude of transitions between two states varying only in transition symbols, which require considerate time to compute and evaluate.

This method proceeds to represent such sets of transitions between two states with (ideally) only a single minterm representing these transitions. We can therefore apply any previously mentioned optimization methods (or any other optimization method) on such modified automata with minterms as their transition symbols to construct their product without the need to compute, for example, Parikh image with every single transition symbol between two states. We can now compute possibly fewer transitions with the acquired minterms instead. Worst case is that the minterms do not reduce any transition symbols and we continue with the same, unchanged original automata. Minterm computation is quick and practically free optimization, which can be used every time an intersection of finite automata is computed.

We apply the minterm computation before we start executing any optimized algorithms introduced here or any others. Instead of putting original automata as the input to opti-

mized algorithms, we compute minterms for such automata and substitute all transitions with the generated minterms. The new input to optimized algorithms are automata with minterms which can be easier to work with, and their intersection can be computed quickly and more precisely. If we need to use the intersection automaton further, not just to resolve the emptiness problem, we simply substitute the minterms in the product with the corresponding original transition symbols back.

Mintermization of our benchmark automata used in our experiments is however not useful as the benchmark automata do not have multiple transitions between two states which could be changed into a minterm. Nevertheless, it is not hard to imagine instances of problems where minterms are essential.

For instance, take regular expressions. If we want to use our abstractions on finite automata representing regular expressions, we have to use mintermization in advance. Otherwise, gained abstractions would be extremely complex to evaluate for even simple regular expressions. Mintermization would allow us to modify such automata into finite automata with manageable number of transitions and transition symbols.

As a future work, we want to try our abstractions in optimizing automata operations in string solving methods such as [1] or [26]. These methods generate regular expressions with rich alphabets and complex transitions with character classes, which would be unsolvable for state language abstractions which have to consider each transition and its symbol.

As an example, imagine a simple regular expression $a[a-z]^*c.^*$ with its finite automaton. It contains numerous transitions for each of the symbols in character classes. Our abstractions would have to compute with each specific transition symbol. However, if we use mintermization on the automaton first, most of the transitions would be simplified into a few minterms: $\{a\}$, $\{c\}$, $\{b, d-z\}$ and $\{\beta\}$ where β represents a set of the remaining symbols of $.^*$ not included in the previous minterms. Instead of complex automata transitions, we now have a finite automaton with only four transition symbols and eleven transitions. Our abstractions can now easily evaluate compatibility of such regular expressions.

Chapter 4

Experiments

The reference implementation[1] of the proposed optimizations, written in Python 3, as well as a complete table of all of our experiments and their results and graphs is publicly accessible on a Codeberg repository[2].

Benchmark with sets of different finite automata used on our benchmark problems are available on a GitHub repository[3]. These finite automata are obtained from runs of regular model checking tool on verification of pointer program and parametric protocols created in [4] based on method of abstract regular model checking from [5]. Such verification runs often execute operations similar to emptiness problem or product construction.

Our experiments cover the benchmark automata from 40 various categories of verification runs. In the benchmark, there are in total 5707 finite automata. However, each category contains similar finite automata recognizing similar languages. The results of our experiments on our benchmark problems for different combinations of finite automata from the same category are nearly identical. Thus, we choose representative finite automata from each category randomly. In total, we have executed more than 300 various experiment runs for combinations of more than 600 finite automata (over 300 pairs of two finite automata from one category). A timeout of 10 minutes for a single test was used.

We test combinations of finite automata from each category to determine the product construction and decide the emptiness of the finite automata intersection. In our experiments, our main objective is to find out how much our optimizations reduce product state space in both our benchmark problems. We want to know what are the pruning capabilities of both our optimizations and whether they are efficient. Further, we want to compare pruning capabilities of length abstraction and Parikh image abstraction to see whether Parikh image pruning capabilities are higher and by how much.

Our abstractions implemented by the reference implementation are not mature enough to properly compete in reduction of time cost of computation yet. Future work includes efficient implementation and further optimizations of our abstractions. Nevertheless, our experiments show our optimizations can sometimes speed up the execution of both benchmark problems.

In this chapter, we present a few experiments which show and compare the pruning capabilities of our optimizations on our benchmark problems. Second, we show what impact have optimizations on our abstractions on computation time benchmark problems.

---

[1]In the reference implementation, we use Z3 as an SMT solver and automata operations are handled by for our purposes modified library Symboliclib.

[2]https://codeberg.org/Adda/optifa

[3]https://github.com/ondrik/automata-benchmarks/tree/master/nfa/non-vtf/armc

4.1 Length Abstraction

In our first experiment, we want to find out what are the pruning capabilities of length abstraction for both our benchmark problems on our benchmark automata. The graph in Figure 4.1 shows a comparison of product state spaces sizes in unoptimized product construction and our optimized algorithm considering length abstraction for emptiness problem. The graph in Figure 4.2 shows a comparison of product state spaces sizes for unoptimized product and product optimized by length abstraction.

Figure 4.1: Emptiness problem.

Figure 4.2: Product construction.

Figure 4.3: Comparison of state space sizes generated by unoptimized product and product optimized by length abstraction for both benchmark problems. Both axes are in symmetrical logarithmic scale[5], x-axis showing the number of states generated by the unoptimized algorithms, y-axis state space sizes of the optimized algorithms.

As we can see, length abstraction successfully prunes state space in some cases. Notice that for some cases, length abstraction can stop product generation immediately on the first product state. When finite automata have large density of final states, they often accept plenty of different word lengths and length abstraction can have problems with finding incompatible abstractions.

The results where length abstraction have difficulties with pruning product space are influenced by our benchmark automata. The combinations of benchmark automata in each category rarely have empty intersections. Therefore, for our next experiment, we want to see whether slight modifications of input automata can highlight the strengths of length abstraction. To further extend the set of benchmark automata for this experiment, to each category, we add finite automata with slight modifications which we combine with original representatives in our experiments. These modifications imitate generation of variations of the same finite automata with different final states (similar to finite automata generated by string solving method from [1]) or little modifications of transitions (removed transitions or changed transition symbols). We want to see whether we can create a state space to prune with our modifications and observe how the length abstraction can notice the difference.

The following graphs show the results of intersection of combinations of the modified benchmark automaton with the original representative from each category for both deciding the emptiness problem and product construction. The graph in Figure 4.4 shows the comparison of product state spaces sizes in unoptimized and our optimized product construction

---

[5]Plot is linear around 0 instead of logarithmic.

for emptiness problem. Sorted in order of increasing product state space size of the unoptimized product. The graph in Figure 4.5 shows the comparison of product state spaces sizes in unoptimized and our optimized product construction for product construction.


Figure 4.4: Emptiness problem.

Figure 4.5: Product construction.

Figure 4.6: Comparison of state space sizes generated by unoptimized product and product optimized by length abstraction of both our benchmark problems. Both axes are in symmetrical logarithmic scale, x-axis showing the state space size of the unoptimized product, y-axis state space size of the optimized product.

We conclude that length abstraction can usually notice the difference between modified and original automata. Length abstraction is able to prune state space more often, and the pruning capabilities of length abstraction are sufficient for these automata. We can see from the graphs that the larger the unoptimized product gets, the higher impact length abstraction has on the product state space size. Product construction optimized by length abstraction generates much smaller products. Length abstraction accomplishes to eliminate state space explosion in most cases.

It is worth mentioning that we have neglected the number of generated states for our lasso automata. Their states are not in the product, but they are required to computation of the product. As we can see in Figure 4.9, even when counting with lasso states, the total number of generated states in the whole process of the product construction is can be lower than the unoptimized product state space size. The larger the automata are, the better results we get. It is understandable that for smaller original automata the overhead of generating lasso automata is significant in comparison with the small generated product state space sizes. However, the larger the original automata get, the lesser the overhead of the number of lasso states is in comparison with the unoptimized product state space.

We can see that even that the overhead of generating lasso automata for length abstraction is necessary, if length abstraction can prune product states, it still pays off: The number of total states generated for either the product or lasso automata is smaller than the number of states in the unoptimized product.

For the rest of our experiments, we use only the original unmodified automata again.

4.1.1 Length Abstraction Optimization without SMT solver

We optimize evaluation of compatibility of length abstractions by substituting SMT solver with solving linear congruence equations. To show how linear congruences speed up the

Figure 4.7: Emptiness problem.

Figure 4.8: Product construction.

Figure 4.9: Comparison of state space sizes in unoptimized product and product optimized by length abstraction with a sum of states generated for both lasso automata. Both axes are in symmetrical logarithmic scale, x-axis showing the number of states in the unoptimized product, y-axis the number of states in the optimized product.

evaluation of compatibility of length abstractions on the original benchmark automata, we present the following experiment.

The Figure 4.10 shows computation time from our benchmark problems pruned by length abstraction with and without SMT solver.


Figure 4.10: Comparison of time consumption of length abstraction evaluated by SMT solver and length abstraction evaluated without SMT solver, combining both benchmark problems. Both axes are in symmetrical logarithmic scale. They show time consumption in seconds: x-axis length abstraction evaluated by SMT solver, y-axis length abstraction evaluated without SMT solver.

We can see that computation of both benchmark problems is faster for length abstractions solved by linear congruences. This optimization improves significantly computation time. Thus, whenever we use length abstraction, we should only evaluate compatibility with linear congruences instead of SMT solver.

Even though our reference implementation of abstractions is not mature enough properly comprete with time cost of basic product ocnstruction, to get a first impression of how our optimized length abstraction compares, we present an experiment n Figure 4.11 showing the difference in time cost for basic product construction and our optimized length abstraction.

Figure 4.11: Comparison of time consumption of unoptimized product construction and length abstraction evaluated without SMT solver, combining both benchmark problems. Both axes are in symmetrical logarithmic scale. They show time consumption in seconds: x-axis unoptimized product construction, y-axis length abstraction evaluated without SMT solver.

As we can see, with length abstraction optimization of removing SMT solver, we can compute both our benchmark problems in time comparable to basic product construction. If parts of product constructin can be pruned, we can even speed up both our benchmark problems.

Out of all experiments with length abstraction, one weakness of length abstraction is clear. The more final states the original automata have, the more difficult it is to optimize product construction using length abstraction. Every final state increases the number of accepted different lengths of automaton. Therefore, with automata where nearly every state is a final state, length abstraction cannot easily determine which product states can be pruned.

4.2 Parikh Image Computation

Length abstraction can sometimes prune product state space significantly, sometimes cannot. We introduced finer abstraction of state languages using Parikh images. Parikh image abstraction computes more precise over-approximation of the state language which would allow us to prune state space more often, even in cases where length abstraction fails. We aim at improving pruning capabilities of our abstractions. In this section, we show experiments with Parikh image abstractions. First, we are interested in pruning capabilities of Parikh image abstraction. Later, we evaluate optimizations of Parikh image abstraction.

we want to find out what are the pruning capabilities of parikh image abstraction in both our benchmark problems on our benchmark automata. in figure 4.14, we can see how parikh image prunes state space.

We can see that Parikh image prunes state space significantly in many cases. Notice that in some cases, Parikh image is able to stop product construction immediately for the initial product state. In total, Parikh image timed out 6 times on emptiness problem and 40 times on product construction. From this we can deduce, that Parikh image is able to decide emptiness problem even for complex automata. For product construction, Parikh image computation often takes significantly longer.

34

Figure 4.12: emptiness problem.

Figure 4.13: product construction.

Figure 4.14: comparison of state space sizes generated by unoptimized and optimized product construction with parikh image abstraction. both axes are in symmetrical logarithmic scale, showing state space sizes: x-axis of unoptimized product, y-axis of optimized product.

In the next experiment, we want to find out whether the Parikh image abstraction has higher pruning capabilities than length abstraction. The Figure 4.15 compares pruning capabilities of length and Parikh image abstractions.


Figure 4.15: Comparison of state spaces generated by length abstraction and Parikh image abstraction, combining both benchmark problems. Both axes are in symmetrical logarithmic scale. Axis show product state space sizes: x-axis for length abstraction, y-axis for Parikh image abstraction.

Clearly, we can conclude from the experiment that Parikh image often optimizes the product state space more than length abstraction (at worst products are equal). Thus, pruning capabilities of Parikh image abstraction are higher than of length abstraction. In many cases, Parikh image optimization is able to prune vast state space by determining incompatible abstractions even if length abstractions are compatible. It is concluded that Parikh image is more precise abstraction, allowing us to prune more aggressively.

Furthermore, notice the dots at the bottom of the graph. Here, Parikh image is able to determine that product language is empty on the first product state and immediately stop the product construction.

4.2.1   Incremental SMT solving

Incremental SMT solving proves to be a great improvement to the Parikh image computation optimization. We want to know how large part of Parikh image formulae can be precomputed and how many conjuncts have to be recomputed for each state. The number of conjuncts in Parikh images depends on the number of states in finite automata, the number of transitions and the number of initial or final states. See Table 4.1 for an example comparison of the number of all conjuncts in Parikh image, conjuncts common to all states (persistent clauses) and state specific conjuncts (state specific clauses).

Product States	All Conjuncts	Persistent Conjuncts	State Specific Conjuncts	Ratio
434	2652	1782	870	67.2%

Table 4.1: An example proportion of persistent and state specific conjuncts in Parikh image computation with incremental SMT solving optimization. *Product States* column shows the number of product states in the whole intersection product, *All Conjuncts* column shows the number of conjuncts in each computed Parikh image, *Persistent Conjuncts* column shows the number of persistent conjuncts in the whole Parikh image (out of the all Parikh image conjuncts), *State Specific Conjuncts* column states how many Parikh image conjuncts have to be recomputed for each product state and *Ratio* column shows the ratio of persistent conjuncts in all Parikh image conjuncts.

In this example, for a product of 434 states, each product state Parikh image contains 2652 conjuncts. From those, 1782 conjuncts are persistent conjuncts and the remaining 870 are state specific conjuncts. A proportional ratio of persistent conjuncts in whole Parikh image is around 67.2%. The number of persistent conjuncts means around 70% of computed Parikh image conjuncts can be precomputed once and used for the whole product generation, and SMT solver can use its cache for efficient evaluation of parts of the Parikh image formulae. Only 30% of conjuncts must be computed repeatedly for each product state.

Even if our abstractions are not mature enough to propely optimize time cost, we want to get a first immpression of what is the cost of more precise pruning capabilities of Parikh image abstraction in both our benchmark problems on our benchmark automata. In Figure 4.18, we can see how Parikh image optimized by incremental solving cost compares to length abstraction optimized by SMT solver substitution with solving linear congruences.

As we can see, the time cost for both problems is higher for Parikh image than for length abstraction. However, time cost of emptiness problem is less affected by the complexity of finite automata. For product construction, Parikh image may take longer to decide compatibility of all product states. We can conclude from the experient that Parikh image time cost of more precise pruning capabilities is present, but in many cases, Parikh image is able to finish the solving of the problem with reasonable time cost.

4.2.2   Precise Timeout Selection

For our benchmark automata, we have experimentally concluded the ideal timeout for SMT solver to solve Parikh image abstractions compatibility is around $600ms$. This gives SMT solver enough time to compute most incompatible cases while it does not wait too long for the confirmation of satisfiability of Parikh image formulae for compatible cases. Our benchmark automata have however large numbers of transitions from each state, and therefore

Figure 4.16: emptiness problem.

Figure 4.17: product construction.

Figure 4.18: Comparison of the time cost of more precise pruning capabilities of Parikh image abstraction with incremental SMT solving and length abstraction optimized by SMT solver substitution with solving linear congruences. Both axes are in symmetrical logarithmic scale, showing time cost in seconds: x-axis of length abstraction, y-axis of Parikh image abstraction.

our timeout might not work best for other types of automata and their complexity. We suggest trying running our optimizations first without any timeout and then, according to the results, adjust the timeout according to the needs of given operations and the complexity of used automata.

4.3   Combination of State Language Abstractions

When we combine length and Parikh image abstraction optimizations in one algorithm, we want to help length abstraction to more precisely prune state space and reduce the number of product states for which Parikh image must be computed. In Figure 4.19, we can see how many product states can be skipped with our skippable states optimization, how many states is pruned by length abstraction and how many states is pruned by Parikh image abstraction (if length abstraction resolves length abstraction formulae as satisfiable). We provide comparison of pruning capabilities of both abstractions on the same automata.

We can clearly see that substantial number of states can be resolved by skippable states optimization or pruned by length abstraction. Parikh images do not have to be computed for any of these states. For the rest of the states, Parikh image have to be computed. We see here again how Parikh image abstraction is precise and helps when length abstraction cannot: Parikh image can prune large numbers of states even though length abstraction fails to prune them. The red and blue column together represent the number of states in the intersection, the yellow and green column the number of pruned states. In total, large parts of product can be pruned. We conclude that combined algorithm using both length abstraction Parikh image abstraction prunes state space really well.

Our abstractions are not mature enough to propely optimize time cost, but to get an impression of what is the time difference in product construction with Parikh image abstraction and with combined algorithm in both our benchmark problems on our benchmark automata, we present the following experiment. In Figure 4.22, we can see how Parikh im-

Figure 4.19: Summary comparison of pruning capabilities of length and Parikh image state language abstractions in optimization algorithm combining both abstractions with optimizations by skippable states. Both axes are in the logarithmic scale: x-axis shows the number of processed product states in product construction (how many product states we have considered in total), y-axis the number of states resolved by respective abstractions. The blue column shows the number of skipped states (we do not need to evaluate abstractions for them), the red column the number of states with compatible both length and Parikh image abstractions, the yellow column the number of states pruned by Parikh image abstraction, but not by length abstraction and the green column the number of states pruned by length abstraction alone. In the graph, we have summed the number of states resolved by each abstraction or optimization according to the number of processed states as follows (from left to right): 0 to 499, 500 to 999, 1000 to 1999 and 2000+

age cost compares to combined approach using both length and Parikh image abstractions.

We can see that length abstraction pruning capabilites can help prune some product state which do not have to be evaluated by Parikh image abstraction and consequently speed up the product construction. There is also a cost of generating lasso automata with our abstractions, which can slightly increase the time cost in cases where there are no product state that can be pruned and both Parikh image and length abstraction have to be evalauted for each one of them.

4.4   Results

The experiments show that our abstractions often prune large parts of the product state space. Length abstraction pruning capabilities are decent, but sometimes it fails to prune state space for intersection of automata with multiple accepted lengths. Pruning capabilities of Parikh image are much higher. Parikh image often succeeds in pruning states where length abstraction fails. Further optimizations of the abstractions have impact on performance of our abstractions. State language abstractions are combinable without affecting pruning capabilities.

Figure 4.20: emptiness problem.

Figure 4.21: product construction.

Figure 4.22: Comparison of the time cost of Parikh image abstraction and combined algorithm using both our abstractions. Both axes are in symmetrical logarithmic scale, showing time cost in seconds: x-axis of Parikh image abstraction, y-axis of combined algorithm.

Chapter 5

Conclusion

The most demanding parts of the intersection computation is the generation of product states and transitions of the product automaton. We tried to reduce the size of the generated state space by omitting the states which cannot lead to any final state by deciding the emptiness of such state languages using various state language abstractions over the finite automata states, such as length abstraction using lasso automata or Parikh image computation based on Parikh's theorem. Each approach has been experimentally tested and further optimizations to the proposed algorithms were introduced.

According to our experiments, product state space can be reduced substantially. Pruning capabilities of our abstractions are satisfactory, and their optimizations have high impact on computation time. We get great results especially for intersections with long lines or for intersections of automata which differ in accepted lengths. Experiments show our algorithm generates smaller state spaces for both resolution of emptiness problem and product construction. Our abstractions are based on over-approximating abstraction of state languages.

We have concluded that length abstraction is fast and coarse abstraction, Parikh image precise but expensive. Our abstractions can be combined, parallelized and further extended.

Due to our discoveries, as a future work, we want to continue working on our state language abstractions, optimize their performance with efficient implementation and explore possibilities of additional improvements of these abstractions. We also want to parallelize evaluation of the compatibility of the abstractions. Further combinations with other abstraction techniques are in consideration, too.

We have not encountered similar approaches to optimization of product construction using length or Parikh image abstractions to compare our results with. Techniques using abstraction were explored especially in a field of program analysis, in the context of automata problem-solving are relevant namely CEGAR [7], IC3/PDR [20, 6, 21, 30, 10] or IMPACT [27]. There were some experiments using techniques similar to length abstraction using information about length constraints [3] to speed up string solving. There are also methods based on the interpolation-based approach of McMillan [2, 16].

It might be worth investing into combining our orthogonal approach with some of the previously mentioned abstraction methods to see how the generated product state space is affected. All the mentioned techniques have proven efficient in hardware or software verification, and they can be applied in automata too.

Bibliography

[1] ABDULLA, P. A., ATIG, M. F., CHEN, Y.-F., HOLÍK, L., REZINE, A. et al. String Constraints for Verification. In: BIERE, A. and BLOEM, R., ed. *Computer Aided Verification*. Cham: Springer International Publishing, 2014, p. 150–166. ISBN 978-3-319-08867-9.

[2] AMLA, N. and MCMILLAN, K. L. Combining Abstraction Refinement and SAT-Based Model Checking. In: *TACAS*. Springer, 2007, vol. 4424, p. 405–419. Lecture Notes in Computer Science.

[3] BERZISH, M., KULCZYNSKI, M., MORA, F., MANEA, F., DAY, J. D. et al. An SMT Solver for Regular Expressions and Linear Arithmetic over String Length. In: SILVA, A. and LEINO, K. R. M., ed. *Computer Aided Verification*. Cham: Springer International Publishing, 2021, p. 289–312. ISBN 978-3-030-81688-9.

[4] BOUAJJANI, A., HABERMEHL, P., HOLÍK, L., TOUILI, T. and VOJNAR, T. Antichain-Based Universality and Inclusion Testing over Nondeterministic Finite Tree Automata. In: IBARRA, O. H. and RAVIKUMAR, B., ed. *Implementation and Applications of Automata*. Berlin, Heidelberg: Springer Berlin Heidelberg, 2008, p. 57–67. ISBN 978-3-540-70844-5.

[5] BOUAJJANI, A., ROGALEWICZ, A., HABERMEHL, P. and VOJNAR, T. Abstract regular tree model checking. In: *ENTCS*. 2006, p. 149(1):37–48.

[6] BRADLEY, A. R. and MANNA, Z. Checking Safety by Inductive Generalization of Counterexamples to Induction. In: *FMCAD*. IEEE Computer Society, 2007, p. 173–180.

[7] CLARKE, E. M., GRUMBERG, O., JHA, S., LU, Y. and VEITH, H. Counterexample-Guided Abstraction Refinement. In: *CAV*. Springer, 2000, vol. 1855, p. 154–169. Lecture Notes in Computer Science.

[8] COLÓN, M. and URIBE, T. E. Generating Finite-State Abstractions of Reactive Systems Using Decision Procedures. In: *CAV*. Springer, 1998, vol. 1427, p. 293–304. Lecture Notes in Computer Science.

[9] CONRAD, K. *Divisibility and Greatest Common Divisors* [online]. [cit. 2021-08-12]. Available at: https://kconrad.math.uconn.edu/blurbs/ugradnumthy/divgcd.pdf.

[10] COX, A. and LEASURE, J. Model Checking Regular Language Constraints. *CoRR*. 2017, abs/1708.09073.

[11] D'ANTONI, L. and VEANES, M. Minimization of symbolic automata. *Proceedings of the 41st ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages.* 2014, p. 541–553.

[12] ESPARZA, J. *Automata Theory: An Algorithmic Approach* [online]. 2017 [cit. 2020-10-31]. Available at: https://www7.in.tum.de/~esparza/automatanotes.html.

[13] ESPARZA, J., GANTY, P., KIEFER, S. and LUTTENBERGER, M. *Parikh's theorem: A simple and direct automaton construction.* June 2011. DOI: 10.1016/j.ipl.2011.03.019.

[14] FIEDOR, T., HOLÍK, L., JANKU, P., LENGÁL, O. and VOJNAR, T. Lazy Automata Techniques for WS1S. In: *TACAS (1).* 2017, vol. 10205, p. 407–425. Lecture Notes in Computer Science.

[15] FIEDOR, T., HOLÍK, L., LENGÁL, O. and VOJNAR, T. Nested antichains for WS1S. *Acta Informatica.* 2019, vol. 56, no. 3, p. 205–228.

[16] GANGE, G., NAVAS, J. A., STUCKEY, P. J., SØNDERGAARD, H. and SCHACHTE, P. Unbounded Model-Checking with Interpolation for Regular Language Constraints. In: *TACAS.* Springer, 2013, vol. 7795, p. 277–291. Lecture Notes in Computer Science.

[17] GANTY, P., MAQUET, N. and RASKIN, J.-F. Fixed point guided abstraction refinement for alternating automata. *Theoretical Computer Science.* 2010, vol. 411, no. 38, p. 3444–3459. DOI: https://doi.org/10.1016/j.tcs.2010.05.037. ISSN 0304-3975. Implementation and Application of Automata (CIAA 2009). Available at: https://www.sciencedirect.com/science/article/pii/S0304397510003294.

[18] GRAF, S. and SAÏDI, H. Construction of Abstract State Graphs with PVS. In: *CAV.* Springer, 1997, vol. 1254, p. 72–83. Lecture Notes in Computer Science.

[19] GREICIUS, A. *Linear Congruences* [online]. 2012 [cit. 2021-08-12]. Available at: http://gauss.math.luc.edu/greicius/Math201/Fall2012/Lectures/linear-congruences.article.pdf.

[20] HODER, K. and BJØRNER, N. Generalized Property Directed Reachability. In: *SAT.* Springer, 2012, vol. 7317, p. 157–171. Lecture Notes in Computer Science.

[21] HOLÍK, L., JANKU, P., LIN, A. W., RÜMMER, P. and VOJNAR, T. String constraints with concatenation and transducers solved efficiently. *Proc. ACM Program. Lang.* 2018, vol. 2, POPL, p. 4:1–4:32.

[22] HOLÍK, L., LENGÁL, O., SÍČ, J., VEANES, M. and VOJNAR, T. Simulation Algorithms for Symbolic Automata. In: *Proc. of 16th International Symposium on Automated Technology for Verification and Analysis.* Springer Verlag, 2018, vol. 11138, no. 1, p. 109–125. DOI: 10.1007/978-3-030-01090-4_7. ISBN 978-3-030-01089-8. Available at: https://www.fit.vut.cz/research/publication/11801.

[23] HOOIMEIJER, P. and VEANES, M. An Evaluation of Automata Algorithms for String Analysis. In: JHALA, R. and SCHMIDT, D., ed. *Verification, Model Checking, and Abstract Interpretation.* Berlin, Heidelberg: Springer Berlin Heidelberg, 2011, p. 248–262. ISBN 978-3-642-18275-4.

[24] JANKŮ, P. and TUROŇOVÁ, L. Solving String Constraints with Approximate Parikh Image. In: MORENO DÍAZ, R., PICHLER, F. and QUESADA ARENCIBIA, A., ed. *Computer Aided Systems Theory – EUROCAST 2019.* Cham: Springer International Publishing, 2020, p. 491–498. ISBN 978-3-030-45093-9.

[25] KOZEN, D. C. Parikh's Theorem. In: *Automata and Computability.* Berlin, Heidelberg: Springer Berlin Heidelberg, 1977, p. 201–205. ISBN 978-3-642-85706-5. Available at: https://doi.org/10.1007/978-3-642-85706-5_35.

[26] LIN, A. W. and BARCELÓ, P. String solving with word equations and transducers: towards a logic for analysing mutation XSS. In: *POPL.* ACM, 2016, p. 123–136.

[27] MCMILLAN, K. L. Lazy Abstraction with Interpolants. In: *CAV.* Springer, 2006, vol. 4144, p. 123–136. Lecture Notes in Computer Science.

[28] SIEGEL, S. F. and YAN, Y. Action-Based Model Checking: Logic, Automata, and Reduction. In: *CAV (2).* Springer, 2020, vol. 12225, p. 77–100. Lecture Notes in Computer Science.

[29] SIPSER, M. *Introduction to the Theory of Computation.* 3rdth ed. Cengage Learning, 2013. ISBN 13: 978-1-133-18779-0.

[30] WANG, H., TSAI, T., LIN, C., YU, F. and JIANG, J. R. String Analysis via Automata Manipulation with Logic Circuit Representation. In: *CAV (1).* Springer, 2016, vol. 9779, p. 241–260. Lecture Notes in Computer Science.

Appendix A

Contents of the Included Storage Media

The following list shows the contents of the included storage media. Listed are only the folders on the highest levels in the folder hierarchy.

- `optifa/`: The main folder with reference implementation of state language abstractions and all related files.

    - `docs/`: The LaTeX source files for this paper.
    - `results/`: The results gained by our experiments.
    - `src/`: The implementation of our optimizations and scripts to run them.
    - `basicDFAs/`: Example finite automata in Timbuk format used in this paper.

- `Symboliclib`: Implementation of the external library Symboliclib with our modifications included.

Appendix B

Reference Implementation Manual

Our reference implementation was tested on GNU/Linux (kernel 5.15.37-1-lts), but it should run on any Unix-like system, possibly even on other operating systems. In order for the reference implementation to work, you need the following programs: Python 3.10[1] or higher, Python library Symboliclib with our modifications and additions[2] and Z3 solver API for Python: Z3Py from Z3 solver repository[3]. Further, to run comparison tests of our optimizations, a command-line benchmarking tool hyperfine[4]. The accepted finite automata file format is Timbuk[5].

Each program can be run with `--help` flag to show a quick help message explaining how to run the program.

To run tests for all our state language abstractions for a specific category (directory with finite automata) or all categories in a directory (for all subdirectories) with `run_tests.py` as follows:

```
./run_tests.py -r <root_directory> -n <experiments_number_per_category> -o <output_file>
```

You can run tests for all our state language abstractions for a specific combination of finite automata with `run_tests.py` as follows:

```
./run_tests.py --single -a <finite_automaton_A> -b <finite_automaton_B> -o <output_file>
```

Separate optimizations can be run with their respective scripts:

- length abstraction with `resolve_satisfiability_length_abstraction.py`, and

- Parikh image abstraction with `resolve_satisfiability_parikh_image.py`.

Combined optimization algorithm using both length and Parikh image abstractions can be run with `resolve_satisfiability_combined.py`.

Each program offers various flags and required or optional arguments to adjust the run according to our requirements: Whether to construct a full product or just test emptiness of the intersection, which abstraction-specific optimizations to enable, where to store results, etc.

Automata with transitions replaced by minterms can be generated with `get_minterms.py`.

---

[1]https://www.python.org/
[2]https://codeberg.org/Adda/symboliclib/; Remember to add Symboliclib to Python path.
[3]https://github.com/Z3Prover/z3; Remember to add Z3Py API to Python path.
[4]https://github.com/sharkdp/hyperfine
[5]https://gitlab.inria.fr/regular-pv/timbuk/timbuk/-/wikis/Specification-File-Format

Appendix C

Complete Optimization Algorithm

Input : NFA $A_1 = (Q_1, \Sigma, \delta_1, I_1, F_1)$, NFA $A_2 = (Q_2, \Sigma, \delta_2, I_2, F_2)$
Output: NFA $P = (A_1 \cap A_2) = (Q, \Sigma, \delta, I, F)$ with $L(A_1 \cap A_2) = L(A_1) \cap L(A_2)$

1 $Q, \delta, F \leftarrow \varnothing$
2 $I \leftarrow I_1 \times I_2$
3 $W \leftarrow I$
4 $res \leftarrow False$
5 $solved \leftarrow \varnothing$
6 addPersistentClauses()
7 **while** $W \neq \varnothing$ **do**
8   **picklast** $[q_1, q_2]$ **from** W
9   **add** $[q_1, q_2]$ **to** $solved$
10   **if** skippable($[q_1, q_2]$) **then**
11 $res \leftarrow True$
12   **else**
13   **if** $\alpha^{LA}(q_1) \wedge \alpha^{LA}(q_2)$ ***is*** unsat **then**
14 $res \leftarrow False$
15   **else**
16   smtSolverPush()
17   addStateSpecificClauses($[q_1, q_2]$)
18 $res \leftarrow \alpha^{PI}(q_1) \wedge \alpha^{PI}(q_2)$ is sat
19   smtSolverPop()
20   **if** $res = Unknown$ **then**
21 $res \leftarrow True$

22   **if** $res = True$ **then**
23   **add** $[q_1, q_2]$ **to** Q
24   **if** $q_1 \in F_1$ ***and*** $q_2 \in F_2$ **then**
25   **add** $[q_1, q_2]$ **to** F
26   **forall** $a \in \Sigma$ **do**
27   **forall** $q_1' \in \delta_1(q_1, a), q_2' \in \delta_2(q_2, a)$ **do**
28   **if** $[q_1', q_2'] \notin solved$ ***and*** $[q_1', q_2'] \notin W$ **then**
29   **add** $[q_1', q_2']$ **to** W
30   **add** $[q_1', q_2']$ **to** $\delta([q_1, q_2], a)$

Algorithm 11: Product construction using both length abstraction and Parikh image computation and all their optimizations.